



NYU

**TANDON SCHOOL
OF ENGINEERING**

Data-Driven Material Selection for Fire-Safe Spacecraft Design

Final Report

Manef Elimami

New York University

August 2026

Abstract

Spacecraft fire safety requires material decisions under atmospheric, flow, geometric, and gravitational conditions for which direct testing is expensive and unevenly distributed. This report consolidates the project’s curated literature database and evaluates whether supervised machine learning can support two tasks: binary prediction of sustained ignition and regression of flame-spread rate (FSR). The evaluation deliberately separates interpolation among conditions represented within known papers from extrapolation to entire held-out papers. After deduplication, the ignition task contains 4,487 observations from 87 canonical papers; the FSR task contains 2,531 usable observations from 66 papers. Preprocessing, hyperparameter search, threshold selection, and augmentation are confined to training partitions.

For ignition, the selected row-level XGBoost model reaches a ROC–AUC of 0.9076 ± 0.0102 , whereas the selected unseen-paper XGBoost model reaches 0.7183 ± 0.0444 . For FSR, the selected random forest reaches $R^2 = 0.8790 \pm 0.0178$ and $\text{RMSE} = 0.0063 \pm 0.0005$ m/s for row-level interpolation; paper-grouped evaluation falls to $R^2 = 0.4576 \pm 0.2215$ and $\text{RMSE} = 0.0122 \pm 0.0041$ m/s. SHAP and permutation analyses identify oxygen, flow and transport quantities, thermal time scales, and fuel properties as important model inputs, but these associations are not interpreted causally. The central engineering result is therefore conditional: the database supports useful within-domain ranking and screening, but the generalization gap and source heterogeneity preclude using the models as stand-alone certification evidence. The models are positioned as traceable decision-support tools for prioritizing tests, identifying poorly covered regimes, and guiding database expansion.

Statement of scope and evidence

All numerical claims in this report are taken from version-controlled database validation files, frozen split metadata, model-selection records, and generated evaluation tables. The oral-defense deck is used to recover the intended narrative, not as the numerical source of record; it explicitly marks one metric page as not updated. Although the defense question mentioned extinction, completed pipelines exist only for ignition and FSR. Consequently, no extinction model or result is claimed here.

Contents

Abstract	i
Statement of scope and evidence	ii
Nomenclature	v
1 Introduction	1
1.1 Spacecraft fire safety	1
1.2 Flammability, quenching, and blowoff	1
2 Literature Review	4
2.1 Review strategy and source roles	4
2.2 Spacecraft fire safety as a material-selection problem	4
2.2.1 Lessons from incidents and evolving practice	4
2.2.2 Qualification tests and their limits	5
2.3 Ignition and flame spread under reduced buoyancy	5
2.3.1 Coupled solid- and gas-phase processes	5
2.3.2 Microgravity transport	5
2.3.3 Atmosphere, pressure, and oxygen	6
2.4 Electrical-wire flame spread	6
2.5 Machine learning for flame-spread prediction	7
2.5.1 Closest prior work	7
2.5.2 Model choice for tabular data	7
2.6 Physics-informed machine learning	8
2.6.1 PINN formulation	8
2.6.2 Relevance and limitation for combustion tables	8
2.7 Synthesis, research gap, and objectives	9
2.7.1 Research gap and data-driven motivation	9
2.7.2 Research question and objectives	9
3 Methodology	10
3.1 Assembly of the microgravity flammability database	10
3.1.1 Literature screening and inclusion	10
3.1.2 Extraction strategy	11
3.1.3 Row structure, outputs, and quality control	11
3.1.4 Source-size and class imbalance	12
3.2 Methods common to both machine-learning pipelines	13
3.2.1 Benchmarking motivation and model families	14
3.2.2 Features, leakage exclusions, and preprocessing	14
3.2.3 Nested evaluation	15
3.2.4 Interpolation versus extrapolation	15
3.2.5 Paper weighting, uncertainty, and refitting	15
3.3 Ignition-classification pipeline	17
3.3.1 Task, probabilities, and metrics	17
3.3.2 Candidate design	17
3.3.3 Pipeline architecture	19
3.3.4 Interpolation protocol	19
3.3.5 Extrapolation and leave-one-paper-out	20
3.4 Flame-spread-rate regression pipeline	20
3.4.1 Task and metrics	20
3.4.2 Candidate design	20
3.4.3 Pipeline architecture	20
3.4.4 Interpolation protocol	22
3.4.5 Extrapolation protocol	22
3.5 Reproducibility and audit trail	22
3.6 Explainability and physical interpretation	23

4	Results	24
4.1	Ignition classification	24
4.1.1	ROC, precision–recall, and the fire-safe operating question	24
4.1.2	Interpolation and extrapolation by model family	25
4.1.3	Ignition explainability	27
4.2	Flame-spread-rate regression	27
4.2.1	Interpolation and extrapolation together	27
4.2.2	Held-out parity and error structure	28
4.2.3	FSR explainability	29
4.3	Synthesis	29
5	Discussion	31
5.1	Limitations	31
5.1.1	Database heterogeneity and the unit of generalization	31
5.1.2	Missingness, measurement comparability, and representation	31
5.1.3	Imbalance and observational sampling	32
5.1.4	Threshold policies and the operating points they induce	32
5.1.5	Model dependence, metrics, and uncertainty	34
5.1.6	SHAP and the boundary between explanation and causation	34
5.2	Comparison with the Literature	34
5.2.1	Rivera/Jose wire-specific machine learning	34
5.2.2	XGBoost and the tabular-learning literature	35
5.2.3	Physics-informed learning and combustion knowledge	35
5.2.4	Spacecraft fire-safety context	36
5.3	Real-World Applications (approximately 30%)	36
5.3.1	Screening and prioritisation	36
5.3.2	Experiment design and the unexplored oxygen–pressure domain	36
5.3.3	Material and design decision support	37
5.3.4	Uncertainty, abstention, and out-of-distribution use	37
5.3.5	No replacement for certification or qualification	38
5.3.6	A staged deployment concept	38
5.4	Future Work (approximately 10%)	38
6	Conclusion	39
A	Technical Appendices	40
A.1	Feature dictionary	40
A.2	Metric definitions	41
A.3	Reproducibility workflow	41
A.4	Traceability matrix	42
A.5	Oral-defense reconciliation	42
A.6	Model portfolios	42

Nomenclature


Symbol or acronym	Definition
c_p	Specific heat capacity, J/(kg K)
g	Acceleration level normalized by terrestrial gravity
k	Thermal conductivity, W/(m K)
α	Thermal diffusivity, m ² /s
ρ	Density, kg/m ³
FSR	Flame-spread rate, m/s
F1	Harmonic mean of precision and recall
FPR	False-positive rate
LOPO	Leave-one-paper-out
MAE	Mean absolute error
MCC	Matthews correlation coefficient
ML	Machine learning
MLP	Multilayer perceptron
OOD	Out of distribution
PR–AUC	Area under the precision–recall curve
RF	Random forest
RMSE	Root-mean-square error
ROC–AUC	Area under the receiver-operating-characteristic curve
SHAP	Shapley additive explanations
TPR	True-positive rate (recall, sensitivity)
XGBoost	Extreme gradient-boosted decision trees

1 Introduction

1.1 Spacecraft fire safety

Fire safety is a system-level requirement for crewed spacecraft. A pressurized vehicle necessarily combines oxidizer, combustible materials, and potential ignition sources in a confined volume where evacuation, replacement of damaged equipment, and atmospheric clean-up may be impossible. Electrical systems are indispensable to vehicle operation, yet faults, overheated components, and degraded insulation place ignition energy next to polymeric fuels. Material selection is consequently a compromise among fire performance, function, mass, manufacturability, and mission architecture rather than a simple elimination of all combustible items [2].

Apollo 1 remains the clearest demonstration of coupled fire hazards. During a ground test on 27 January 1967, a fire developed rapidly in the command module's high-pressure pure-oxygen atmosphere. The exact ignition source was not established by the investigation, but electrical wiring was considered a likely source; extensive combustible material, the atmosphere, and the inward-opening hatch made the event catastrophic. The post-fire views in figure 1.1 show why specimen flammability cannot be separated from installation and vehicle configuration. The accident led to atmospheric, materials, wiring, hatch, and test-program changes, including extensive component and full-scale mock-up testing [2].



(a) Apollo 204 command module during the post-fire investigation.



(b) Prime crew: Grissom, White, and Chaffee.

Figure 1.1: Apollo 1 command module and crew. NASA public-domain photographs [6, 7].

Deep-space missions increase the consequences of such an event. A distant vehicle or surface habitat cannot assume rapid evacuation, immediate resupply, or prompt replacement of discharged extinguishers and damaged life-support hardware. Proposed low-pressure atmospheres, elevated oxygen fraction, lunar or Martian gravity, radiation aging, dust, waste accumulation, and long autonomous operation all interact with fire risk [2]. A material qualified in one terrestrial configuration therefore cannot be assumed to keep the same hazard ranking in every spacecraft environment.

1.2 Flammability, quenching, and blowoff

In this report, *flammability* is not reduced to a single binary property. It is represented by five complementary experimental observables:

These variables describe different stages and consequences of a fire. Ignition addresses initiation; flame length, FSR, and HRR characterize growth; smoke and aerosols affect detection and habitability. Their availability is uneven across the literature. The implemented machine-learning pipelines model only ignition and FSR. Flame length, HRR, and smoke/aerosols are retained as database outputs but excluded from predictors to prevent outcome leakage.

The stability of an established flame is bounded by two distinct extinction mechanisms. At high imposed flow or stretch, the gas residence time through the reaction zone becomes too short relative to the chemical reaction time. The flame detaches or is extinguished by *blowoff*, conventionally associated with a low Damköhler number. Self-induced buoyant flow can also produce blowoff during upward spread in normal gravity [3].

Table 1.1: Five variables used to describe material flammability in the database.

Variable		Engineering interpretation
Ignition		Whether the imposed energy and environment establish a flame.
Flame length		Spatial extent of the visible flame under the test condition.
Flame-spread (FSR)	rate	Speed of the advancing flame front, reported in m/s.
Heat-release (HRR)	rate	Rate of chemical energy release, reported in W when available.
Smoke/aerosols		Production of particulate and condensable combustion products relevant to visibility, toxicity, and detection.

At the opposite limit, very weak flow supplies insufficient convective heat feedback and oxygen transport to offset radiative and conductive losses. The flame shrinks and extinguishes by *quenching*. The low-speed quenching branch is experimentally accessible in microgravity because buoyancy no longer imposes a minimum natural-convection velocity. In ordinary gravity, buoyant flow masks this near-zero-flow branch; this is the precise sense in which the quenching limit considered here is observed only under microgravity conditions, not a claim that every possible form of flame quenching is exclusive to space.

Plotting limiting oxygen concentration against forced stretch or flow gives the U-shaped boundary in figure 1.2. The left branch is the microgravity low-flow quenching limit, the right branch is the high-flow blowoff limit, and stable burning occupies the region above the boundary. The minimum identifies the lowest oxygen concentration at which any tested flow can sustain combustion. This map demonstrates why oxygen concentration alone is insufficient: the same material and atmosphere may be non-flammable at very low and very high flow yet burn at an intermediate flow.

The remainder of this report reviews the spacecraft, combustion, and machine-learning literature that frames the research gap; states the resulting question and objectives; describes the database and evaluation design; and reports ignition and flame-spread predictions under interpolation and paper-held-out extrapolation.




Figure 1.2: Schematic U-shaped flammability boundary. The low-velocity branch is quenching and the high-velocity branch is blowoff; curves show the reported fuel-width dependence. Reproduced from Johnston et al. [3].

2 Literature Review

2.1 Review strategy and source roles

We synthesize five sources that address different layers of the research problem. Guibaud et al. provide the spacecraft fire-safety context, historical evidence, and deep-space challenges [2]. Rivera et al. review the physics of flame spread over electrical wires and present the closest prior machine-learning study to the present FSR task [9]. Schwartz-Ziv and Armon provide comparative evidence about model selection for tabular data [11]. Raissi, Perdikaris, and Karniadakis establish the PINN formulation [8], while Karniadakis et al. situate physics-informed learning within the wider scientific-machine-learning field [4].

These source roles must remain distinct. Historical incident counts, spacecraft atmosphere descriptions, and deep-space concerns are attributed to Guibaud et al. Physical and machine-learning statements specific to wire flame spread are attributed to Rivera et al. Claims about tabular algorithm comparisons are limited to the datasets and protocol of Schwartz-Ziv and Armon. PINN equations describe methods reported in the PINN literature; they do not imply that this project’s tabular pipelines solved a combustion PDE. Finally, none of the performance values reported for earlier studies is presented as a result of this project.

Table 2.1: Authoritative sources and their use in this report.

Source	Principal evidence used	Limitation for the present question
Guibaud et al. [2]	Spacecraft incidents, fire-safety strategy, reduced-buoyancy implications, and deep-space challenges	Review and risk context; not a trained ignition or FSR predictor
Rivera et al. [9]	Wire-flame physics, heterogeneous experimental predictors, categorization, augmentation, and ML FSR benchmarks	Focuses on wire FSR; random cross-validation is not strict held-source-out transfer
Shwartz-Ziv and Armon [11]	Controlled comparison of XGBoost with deep models on eleven tabular datasets	General tabular evidence, not combustion-specific evidence
Raissi et al. [8]	Differential-equation residuals, collocation points, and forward/inverse PINN formulation	Assumes suitable governing equations and coordinate-resolved information
Karniadakis et al. [4]	Broader synthesis of physics-informed machine learning and data-physics integration	Research framework rather than a spacecraft-material qualification method

2.2 Spacecraft fire safety as a material-selection problem

2.2.1 Lessons from incidents and evolving practice

The history reviewed by Guibaud et al. shows that spacecraft fire safety evolved through incident investigation, atmospheric redesign, material screening, detection hardware, suppression equipment, and crew procedures. Their chronology identifies twelve acknowledged fire or smoke events. Apollo 1 remains the clearest demonstration of coupled design hazards: an oxygen-rich cabin, combustible interior materials, possible electrical ignition, rapid pressure rise, and an inward-opening hatch combined in a fatal accident during a ground test. The precise ignition source was not established, but subsequent actions included atmospheric changes, removal or protection of combustibles, and extensive component and mock-up flammability testing [2].

Apollo 1 also demonstrated that material behaviour cannot always be inferred from an isolated specimen. Guibaud et al. report that components showing acceptable individual behaviour did not necessarily prevent excessive propagation after integration into a full spacecraft mock-up. The configuration changed the available fuel paths, heat feedback, and interactions among adjacent items. This observation remains relevant to data-driven screening: a row-level material prediction can inform a decision, but cannot by itself represent every installation detail in a complete vehicle [2].

Later incidents shifted attention toward electrical equipment and operational response. The Salyut and Shuttle records include failed fans, short circuits, overheated resistors, fused wires, and failed electronic components. In the five Shuttle events summarized by Guibaud et al., the crew interrupted power before a full-scale fire developed, while onboard detection systems did not provide the decisive first warning. The 1997 Mir fire showed a different failure mode: a lithium-perchlorate oxygen canister produced an energetic flame and dense smoke, obstructed access to one escape vehicle, and could not be directly suppressed by the available water-based foam. Crew members

had to improvise while contending with visibility, equipment access, and respirator problems [2]. The literature’s point is not that every spacecraft fire resembles a wire fire, but that ignition source, material, atmosphere, access, and mitigation interact at system scale.

Present practice uses multiple defence layers. Guibaud et al. describe fire-resistant material selection, configuration analysis, redundant or distributed smoke detection, electrical and ventilation shutdown, portable extinguishers, atmospheric monitoring, crew masks, and rehearsed procedures on the International Space Station. This layered architecture has benefited from earlier programmes, yet it still relies on knowledge of material response and combustion signatures in low gravity. The review argues that uncertainties remain in translating terrestrial test performance to reduced gravity and in demonstrating suppression performance for every relevant scenario [2].

2.2.2 Qualification tests and their limits

Spacecraft material control is intentionally conservative. The test history summarized by Guibaud et al. includes NASA’s vertical flame-propagation screening, in which multiple specimens are exposed to an ignition source under a specified atmosphere. The cited criterion considers both propagation distance and ignition by burning debris. At the vehicle level, a material that does not pass an isolated test may still require configuration analysis when mission-critical function prevents simple removal. Separation, containment, barriers, and elimination of propagation paths become part of the safety argument [2].

This approach is essential but does not produce a continuous map over all operating conditions. A pass/fail test answers a deliberately bounded qualification question. It does not directly estimate a probability of ignition for every pressure and oxygen concentration, nor does it yield FSR over arbitrary flow and gravity. Furthermore, Guibaud et al. cite reduced-gravity observations in which limiting behaviour differs from normal-gravity expectations and describe partial-gravity studies where relative material ranking can change [2]. A data-driven surrogate is useful precisely in this space between discrete experiments: it can organize existing observations, interpolate under controlled assumptions, and expose regions where direct testing remains necessary. It cannot supersede the qualification test from which the safety requirement is defined.

2.3 Ignition and flame spread under reduced buoyancy

2.3.1 Coupled solid- and gas-phase processes

Ignition is a transition rather than an intrinsic label attached permanently to a material. Energy supplied by an igniter or fault heats the solid; heat is stored and conducted through the specimen and, if the material decomposes, volatile fuel is generated. Gas-phase transport then mixes fuel and oxidizer, while losses to the core, surroundings, and unheated material compete with heat feedback. The literature reviewed by Rivera et al. emphasizes that even reduced analytical descriptions require material properties, dimensions, gas properties, and heat-flux estimates that are not uniformly reported across experiments [9]. Consequently, an observed yes/no ignition outcome is conditional on the entire experimental configuration.

A useful scale for solid heating is the thermal diffusion time

$$t_{\text{diff}} \sim \frac{\ell_c^2}{\alpha_s}, \quad \alpha_s = \frac{k_s}{\rho_s c_{p,s}}, \quad (2.1)$$

where ℓ_c is a characteristic solid length, k_s is thermal conductivity in $\text{W m}^{-1} \text{K}^{-1}$, ρ_s is density in kg m^{-3} , $c_{p,s}$ is specific heat in $\text{J kg}^{-1} \text{K}^{-1}$, and α_s is thermal diffusivity in $\text{m}^2 \text{s}^{-1}$. This scaling does not determine ignition by itself. It shows why thickness and thermophysical properties jointly affect the time over which a specimen responds to heating, and why separately recorded database columns can have interacting rather than additive effects.

After ignition, flame spread requires successive heating and pyrolysis of virgin material. Rivera et al. reproduce an analytical opposed-flow flame-spread relation in which spread speed depends on heated length, solid thermal inertia, pyrolysis and ambient temperatures, several flame and external heat-flux terms, radiative loss, and a kinetic contribution [9]. At the level needed here, that balance can be written schematically as

$$v_{\text{FSR}} \sim \frac{\dot{q}_{\text{net}}'' \ell_h}{\rho_s c_{p,s} \delta (T_p - T_0)}, \quad (2.2)$$

where $\dot{q}_{\text{net}}''$ denotes net heating of the unburned solid, ℓ_h is a heated length, δ is a solid thickness scale, T_p is a pyrolysis-temperature scale, and T_0 is the initial temperature. Equation (2.2) is a qualitative energy-balance representation, not the regression equation used by this project. It makes the key coupling explicit: any variable that changes heat feedback, heat storage, or heat loss can alter spread speed.

2.3.2 Microgravity transport

On Earth, gravity-induced buoyancy contributes a characteristic velocity and organizes the flame and smoke plume. In microgravity that contribution is suppressed, so imposed ventilation and diffusion control transport to a much

greater degree. Guibaud et al. connect this change to flammability, spread, smoke movement, detector response, and suppression uncertainty. A fire may produce combustion products that take longer to reach a sensor, while a relatively small ventilation flow can reshape a non-buoyant flame [2]. This is why “zero gravity” cannot be represented only as one more categorical label while all other transport variables are ignored.

Dimensionless groups help organize the competing mechanisms. For characteristic velocity U , length L , gas kinematic viscosity ν , and gas thermal diffusivity α_g ,

$$\text{Re} = \frac{UL}{\nu}, \quad \text{Pe} = \frac{UL}{\alpha_g}, \quad \text{Pr} = \frac{\nu}{\alpha_g}. \quad (2.3)$$

These groups compare inertial to viscous transport, advection to thermal diffusion, and momentum to thermal diffusivity, respectively. Their value in a heterogeneous database is representational: experiments with different dimensions and velocities may occupy related transport regimes. They do not remove the need to preserve dimensional predictors or account for geometry.

Opposed and concurrent flow are physically different configurations. In opposed spread, gas flow is directed toward the advancing flame front; in concurrent spread, flow and flame advance share a direction. Rivera et al. identify spread direction as an important categorical descriptor and report that adding such regime information substantially improved predictions in their compiled wire dataset [9]. The result supports explicit representation of configuration, but it must be interpreted within that study’s cross-validation protocol. It does not prove that a direction label captures all flow physics.

Partial gravity adds further complexity. Guibaud et al. review experiments and modelling that indicate non-monotonic responses for some fire quantities between microgravity and terrestrial gravity. They report different trends for upward and downward propagation and note that material ranking may change between normal, lunar, and Martian gravity [2]. The engineering consequence is that a predictor should not silently infer partial-gravity behaviour from an arbitrary linear interpolation between microgravity and terrestrial gravity. Sparse coverage must instead be treated as an extrapolation problem and reported as such.

2.3.3 Atmosphere, pressure, and oxygen

Oxygen concentration and pressure appear repeatedly in the spacecraft literature because mission atmosphere is both a life-support and structural design choice. Guibaud et al. describe the pure-oxygen, low-pressure atmosphere used in early American flight and the much higher pressure pure-oxygen condition present during the Apollo 1 ground test. Later spacecraft generally adopted nitrogen-diluted atmospheres, while proposed exploration vehicles may again use reduced total pressure and elevated oxygen fraction to address structural mass and extravehicular-activity considerations [2]. Fire response must therefore be considered over a domain, not only at terrestrial sea-level conditions.

Oxygen fraction is not equivalent to oxygen partial pressure, and neither alone captures all pressure effects. For an ideal mixture the oxygen partial pressure is

$$p_{\text{O}_2} = X_{\text{O}_2} P, \quad (2.4)$$

with mole fraction X_{O_2} and total pressure P in pascals. Equal values of p_{O_2} can still coexist with different total gas density and transport properties when diluent and total pressure differ. Rivera et al. consequently include oxygen concentration and pressure as separate predictors in the wire FSR problem [9]. This supports a multivariate representation rather than collapse to a single atmosphere variable.

2.4 Electrical-wire flame spread

Electrical wires are a particularly relevant material system because conductor and insulation form a coupled thermal geometry. Rivera et al. synthesize experiments using laboratory wires with metal cores and polymer insulation. Their review identifies core material and diameter, insulation material and thickness, oxygen concentration, pressure, forced flow, gravity, and orientation as influential quantities [9]. These variables also illustrate why the broader database is heterogeneous: a nominal material may appear in different specimen forms, and the same environment may produce different outcomes for different length scales.

The central conductor has a dual role. A high-conductivity core can move heat ahead of the flame and assist preheating. The same core can drain energy away from the reaction and pyrolysis region, especially when its cross-section is large relative to the insulation. Rivera et al. attribute this balance to core and insulation dimensions and properties, environmental conditions, and gravity [9]. A monotonic rule such as “higher core conductivity always increases FSR” would therefore contradict the reviewed physics. Flexible nonlinear models are attractive because they can represent conditional interactions, but flexibility also increases the need for careful validation.

Insulation behaviour introduces discrete regimes. Polymer type affects thermophysical response, while melting and dripping can change fuel geometry and energy transport. Orientation determines whether gravity assists movement of molten insulation in normal or partial gravity. Rivera et al. therefore encode core material, insulation material, spread direction, and configuration as categorical predictors in addition to numerical properties [9]. This

is an important methodological lesson: a table containing only bulk conductivity, density, and heat capacity may omit morphological or chemical distinctions that those properties do not fully identify.

The wire literature also exposes a transferability issue. Studies may concentrate on one polymer, core, diameter range, apparatus, pressure regime, or flow direction. When observations from the same study are divided randomly among training and test folds, the test set can remain close to the training distribution. Rivera et al. acknowledge that their random ten-fold cross-validation measures robustness within the compiled literature dataset rather than strict transfer to an entirely new source [9]. For material selection, this distinction is critical because a new candidate is often associated with a new material family or experimental campaign.

2.5 Machine learning for flame-spread prediction

2.5.1 Closest prior work

Rivera et al. provide the closest direct precedent for data-driven FSR prediction. They compile approximately 1200 wire-flame observations and describe numerical features for core, insulation, geometry, atmosphere, flow, and gravity, together with categorical indicators for material and spread configuration. They compare linear regression, decision tree, k -nearest neighbours, elastic net, stochastic gradient descent, Bayesian ridge, gradient boosting, support-vector regression, kernel ridge, and a multilayer perceptron [9]. The breadth of this benchmark is useful because it avoids treating one model family as the default.

Their baseline using numerical predictors alone produced weaker performance than versions that included one-hot encoded categorical variables. Under categorized random ten-fold cross-validation, several nonlinear models attained mean coefficients of determination around or above 0.90; the exact values vary by model and augmentation strategy [9]. These are literature results. They cannot be transferred directly to this project’s database, targets, partitions, or code, and they do not establish held-paper generalization.

Rivera et al. also compare bootstrap resampling, small numerical perturbations, and a generative adversarial network (GAN). Bootstrap augmentation was the most consistently beneficial for tree- and neighbourhood-based models in their reported experiments. The authors stress that resampling preserves observed combinations, whereas noise and GAN generation can produce statistically plausible combinations without guaranteeing geometric consistency, conservation, or other physical constraints [9]. This qualification matters: synthetic tabular rows are not new experiments. If material density, conductivity, dimensions, and categorical identity are perturbed independently, the resulting row may not correspond to a realizable specimen.

The prior study reports a small number of negative FSR predictions from unconstrained regressors and correctly identifies them as physically inadmissible [9]. This motivates non-negative output handling, target transformation, or explicit post-fit diagnostics. More generally, machine-learning performance must be judged against physical admissibility as well as aggregate error. A model can obtain a favourable mean score while failing in the low-FSR region that is most relevant to a safety boundary.

2.5.2 Model choice for tabular data

The structure of the combustion database resembles the broad statistical category studied by Schwartz-Ziv and Armon: fixed columns, heterogeneous units, mixed numerical and categorical features, missingness, and no spatial or sequential locality comparable to pixels or words. Their controlled study evaluates four deep tabular architectures and XGBoost across eleven classification and regression datasets, with a common hyperparameter-optimization protocol. XGBoost outperformed the deep models on most of the considered datasets, particularly those not used in the original deep-model papers, and its tuning converged more readily in their experiments [11].

The conclusion is narrower than “deep learning never works for tables.” Schwartz-Ziv and Armon explicitly invoke the no-free-lunch principle and report that an ensemble containing both deep models and XGBoost achieved the best average relative performance in their comparison [11]. For this project, the defensible inference is to include gradient-boosted trees as a strong baseline and compare them fairly with other candidates. Model selection must follow held-out evidence under the relevant engineering split, not reputation or architectural novelty.

Tree ensembles are plausible for this task because thresholded partitions can capture nonlinear interactions and accommodate regime changes, while neural and kernel models offer different smoothness and representation biases. No family removes the need for leakage control. Imputation, scaling, categorical encoding, feature selection, and hyperparameter optimization must be fitted using training data only. Repeated measurements and paper-specific clusters should remain grouped where the evaluation question concerns transfer to a new source. These are methodological requirements of the present project, not performance claims borrowed from the cited tabular benchmark.

2.6 Physics-informed machine learning

2.6.1 PINN formulation

Raissi et al. define PINNs for systems represented by nonlinear differential equations. For a latent field $u(t, \mathbf{x})$,

$$\frac{\partial u}{\partial t} + \mathcal{N}[u; \boldsymbol{\lambda}] = 0, \quad \mathbf{x} \in \Omega, \quad t \in [0, T], \quad (2.5)$$

the neural network approximates u , and automatic differentiation evaluates the residual

$$r_\theta(t, \mathbf{x}) = \frac{\partial u_\theta}{\partial t} + \mathcal{N}[u_\theta; \boldsymbol{\lambda}]. \quad (2.6)$$

Training combines mismatch to available observations and initial or boundary data with residual error at collocation points,

$$\mathcal{L}(\theta) = \frac{1}{N_u} \sum_{i=1}^{N_u} |u_\theta(t_i, \mathbf{x}_i) - u_i|^2 + \frac{1}{N_r} \sum_{j=1}^{N_r} |r_\theta(t_j, \mathbf{x}_j)|^2. \quad (2.7)$$

The framework supports forward inference with known physical parameters and inverse identification in which parameters in $\mathcal{N}$ are learned from observations [8].




Figure 2.1: Example physics-informed neural-network solution from Raissi et al.: sparse initial and boundary observations and interior collocation points constrain the nonlinear Schrödinger solution; predicted profiles (red dashed) are compared with the exact solution (blue). Reproduced from Fig. 1 of [8].

Raissi et al. demonstrate the formulation on nonlinear Schrödinger, Allen–Cahn, Navier–Stokes, Korteweg–de Vries, and Burgers equations, using continuous- and discrete-time variants. They emphasize that PINNs complement rather than replace mature numerical methods and identify unresolved issues involving architecture, optimization, convergence, noise, and uncertainty [8]. Karniadakis et al. broaden the discussion to scientific settings where observational data may be scarce, noisy, or incomplete and prior physical knowledge can constrain an otherwise large hypothesis space [4].

2.6.2 Relevance and limitation for combustion tables

The appeal to spacecraft combustion is evident. Reduced-gravity experiments are costly, and conservation or transport relations contain information not present in a small set of labels. A valid physics-informed constraint could discourage impossible predictions and improve data efficiency. Dimensionless groups, positivity, energy-balance trends, and known symmetries are possible levels of physical structure, ranging from feature engineering to hard or soft constraints. The PINN literature provides a principled reason to investigate such structure [4, 8].

However, the standard PINN formulation requires more than a list of experiment-level rows. It needs a specified state field, spatial and temporal coordinates, a governing differential operator, and sufficient initial or boundary information for the target problem. The curated database records outcomes such as ignition and FSR under varied apparatuses; it is not a common mesh of temperature, species concentration, velocity, and reaction-rate

fields. Applying Equation (2.7) without a defensible operator and boundary conditions would create the appearance of physics without its mathematical content.

This limitation is reinforced by Rivera et al., who conclude that the theoretical framework for their heterogeneous wire data remains incomplete for direct use of fully physics-informed models [9]. The scientifically appropriate hierarchy is therefore: preserve SI units and physically meaningful variables; encode known regimes; enforce indisputable output constraints such as non-negative FSR; evaluate transfer across sources; and reserve PDE-residual PINNs for datasets and subproblems where the governing fields and boundaries are actually available. This project can be physics-aware without claiming a PINN where none was implemented.

2.7 Synthesis, research gap, and objectives

The literature converges on four findings. First, spacecraft fire is a credible system hazard whose consequence increases when evacuation, resupply, and recovery are restricted [2]. Second, ignition and spread depend jointly on material, geometry, atmosphere, flow, and gravity; reduced buoyancy changes transport rather than merely slowing or accelerating a terrestrial flame by a fixed factor [2, 9]. Third, nonlinear supervised models can recover useful relationships from heterogeneous wire-flame tables, especially when physically meaningful categories are retained [9]. Fourth, strong tabular baselines and physical constraints should be evaluated empirically and applied only within the assumptions their source literature supports [4, 8, 11].

2.7.1 Research gap and data-driven motivation

Reduced gravity changes the balance among diffusion, imposed ventilation, radiation, and solid conduction. Experiments are correspondingly expensive, and the published evidence is fragmented across materials, geometries, facilities, atmospheres, and reporting conventions. A finite qualification matrix cannot dense-sample this space, while a high-fidelity reacting-flow model requires boundary conditions and material properties that are often absent from historical reports.

The research gap is therefore not simply a lack of another predictive algorithm. Three specific absences remain. The *target gap* is that the closest machine-learning study concentrates on FSR, whereas an engineering screening workflow also needs an explicit ignition model. The *validation gap* is that randomly splitting rows can place nearly repeated observations from one rig on both sides of the evaluation boundary, creating optimistic evidence for a material-selection question that usually concerns a new campaign. Random cross-validation within a compiled literature table does not measure transfer to a held-out paper, apparatus, or condition regime. The *integration gap* is that broad spacecraft studies establish operational relevance while wire and tabular-learning studies establish local predictive feasibility, but the evidence has not yet been organized into one traceable, leakage-resistant evaluation of both ignition and FSR.

A data-driven response is motivated by this fragmentation rather than by a claim that combustion physics is obsolete. What is missing is a unified database and an evaluation that asks both whether models interpolate among conditions from known campaigns and whether they transfer to a completely held-out paper. That distinction is the difference between recovering nearby points in a familiar apparatus and supporting a screening decision for a new material, geometry, or laboratory.

2.7.2 Research question and objectives

To what extent can supervised models trained on a curated, multi-source microgravity-combustion database predict (i) whether ignition occurs and (ii) the flame-spread rate under spacecraft-relevant material, geometric, atmospheric, flow, and gravity conditions, while maintaining performance when entire experimental papers are withheld?

The work has four objectives:

1. assemble a source-traceable database with consistent variables and SI units from heterogeneous publications;
2. construct clear, reproducible classification and regression pipelines;
3. compare row-level interpolation with paper-grouped extrapolation; and
4. interpret model behavior using held-out diagnostics, permutation importance, and SHAP without treating predictive associations as causal laws.

The intended role is screening and experiment prioritization, not replacement of qualification tests or combustion physics. Only ignition and FSR are modeled. Extinction, flame length, HRR, and smoke/aerosol prediction remain future modeling tasks. This boundary preserves a clear distinction between the database’s definition of flammability and the narrower scope of the validated machine-learning evidence. The project predicts recorded ignition and FSR outcomes; it does not certify flight materials or solve the full reacting-flow problem. A reliable material-selection tool must make clear both what has been learned from existing tests and where a new experiment remains the only defensible next step.

3 Methodology

This chapter specifies the complete chain from literature evidence to held-out predictions. The unit recorded in the database is an experimental condition, but the unit of scientific transfer is a publication: rows from one paper share apparatus, materials, measurement practice, and often a common experimental campaign. The study design therefore treats row-level interpolation and paper-level extrapolation as distinct estimands, confines every data-adaptive operation to training partitions, and preserves a machine-readable audit trail from each reported score to its source rows, split assignment, candidate configuration, and fitted model. These principles were fixed before comparing candidate performance.

3.1 Assembly of the microgravity flammability database

3.1.1 Literature screening and inclusion

The database was assembled before either machine-learning pipeline was designed. Approximately 130 experimental papers and technical reports published from 1966 onward were screened for solid-fuel combustion under microgravity, partial-gravity, or spacecraft-relevant atmospheric and flow conditions. Only about 87 publications exposed sufficiently resolved observations to map ignition or flame-spread rate (FSR) into individual database rows. A paper was useful for modeling only when the experiment, specimen, environment, and measured outcome could be connected at row level; a qualitative conclusion without recoverable conditions could support the literature review but not supervised learning.

The resulting corpus is deliberately broader than one material family or rig. It includes wires, flat specimens, cylinders, and spheres tested on the ground, in drop towers, parabolic aircraft, and spacecraft. Each row retains a full human-readable citation, authors, and DOI or report identifier where available. That provenance is essential because observations from one publication are clustered and because ambiguous values must be checked against the source rather than treated as anonymous numbers.

Dataset Overview - Microgravity Flammability Database (reduced)

Total data points 5,057	Pressure range (kPa) 3.2 - 6343
Number of papers 87	Gravity range (g/g_earth) 0.00 - 4.51
Median points / paper 32	Fuel density range (kg/m3) 25 - 2200
Ignition (Yes) 3,875 (76.6%)	Fuel conductivity range (W/mK) 0.035 - 5.00
Extinction / No-ignition (No) 1,182 (23.4%)	Points with numeric FSR 2,590
O2 mole-fraction range 0.002 - 1.000	FSR range (mm/s) -0.01 - 211

Figure 3.1: Overview of the assembled microgravity flammability table before task-specific cleaning. The corpus contains 5,057 rows from 87 papers (median 32 rows per paper). Ignition is the majority outcome (76.6%). Numeric FSR is available for 2,590 raw rows. Operating envelopes span oxygen mole fraction 0.002–1, pressure from a few kilopascals to several megapascals, gravity from microgravity to several g , and a wide fuel-property range. Modeling later uses SI flame-spread units of m/s; the overview card reports the exploratory FSR range in millimetres per second.

The overview is not a balanced factorial design. A median of 32 rows per paper coexists with campaigns that contribute several hundred nearby conditions from a single rig. That source-size structure is why a random row split is optimistic: conditions from a large paper can appear on both sides of an interpolation split, whereas holding out that paper removes a large, internally correlated block. The same card also shows why two supervised tasks,

not five, are feasible: ignition labels are essentially complete, numeric FSR covers about half the table, and the remaining flammability outputs are too sparse to train.

3.1.2 Extraction strategy

The first attempted extraction workflow was a conventional optical character recognition (OCR) pipeline. It was not viable for this corpus. The sources range from old scanned papers to long NASA memoranda; notation, units, table structures, and page layouts differ; scans and plotted markers can be blurred; and many results appear only in graphs rather than machine-readable text. Generic OCR also confuses subscripts and symbols, separates values from their headers, and cannot reliably infer which experimental condition belongs to a plotted point. Applying it at scale would have created a large table with an unacceptably hidden transcription error rate.

The adopted workflow was therefore AI-assisted but human-verified:

1. the source was read to identify the experiment, sample, controlled conditions, reported outputs, and units;
2. an AI extraction pass was used to recover as much structured information as possible for each experimental row;
3. every extracted row was then checked against the paper by the database curator, and missing or incorrectly associated values were completed or corrected manually;
4. units and categorical names were standardized while the original citation and explanatory notes were retained for traceability.

AI was used as an extraction accelerator, not as an authority. This distinction is important for scientific data: a fluent transcription can still attach a correct number to the wrong series, silently flatten a logarithmic axis, or invent precision that the graph does not contain. Manual review remained the acceptance step. Values read from plots were stored only to the precision supported by marker size and axis resolution; visually inferred trailing digits were not treated as measured information. Ambiguous series assignments or units were resolved from captions and surrounding text, or left missing when the source did not support a defensible interpretation.

For dense plots, the points were digitized with WebPlotDigitizer [10]. Axis calibration was set from the published ticks, including explicit log-axis selection where required; each color or marker series was traced separately; overlapping points were inspected at high zoom; and the exported coordinates were checked against axis limits and values quoted in the text. This procedure was preferred to AI-internal OCR because exact point coordinates, logarithmic scaling, and overlapping markers were otherwise unreliable.

3.1.3 Row structure, outputs, and quality control

The final schema groups fields into citation, specimen, flow, apparatus, outputs, and notes. Inputs cover geometry, dimensions, fuel and core thermophysical properties, oxygen concentration, diluent, gas properties, pressure, flow, dimensionless transport groups, gravity, facility, and ignition conditions. Flammability is documented through the five outputs introduced in table 1.1: ignition, flame length, FSR, heat-release rate, and smoke/aerosols. Not every paper reports every output.

The raw table uses a category banner as its first row and field names as its second. Values were standardized to SI units wherever the source supported a defensible conversion: length in metres, pressure in pascals, velocity in m/s, power in watts, time in seconds, density in kg/m³, conductivity in W/(m K), heat capacity in J/(kg K), and FSR in m/s. Values whose physical dimension did not match a target field were not coerced merely to fill a cell.

Table 3.1: Complete field dictionary of the Excel database (`Microgravity_Database_converted.xlsx`). Completeness is the fraction of the 5,057 assembled rows with a non-empty cell. Modeling roles state whether the field is a predictor, a target, a grouping key, or is excluded from supervised inputs.

Group	Excel field	Units	Compl. (%)	Modeling role
Citation	Article (MLA)	—	100.0	paper identity / grouping
Citation	Authors	—	100.0	provenance; excluded predictor
Citation	DOI	—	93.2	paper identity / grouping
Sample	Geometry of Sample	categorical	100.0	predictor
Sample	Dimensions of sample	m (geometry-specific)	100.0	parsed into length scales
Sample	<code>half_thickness</code>	m	100.0	predictor
Sample	<code>characteristic_length_m</code>	m	100.0	predictor
Sample	Material of sample	free text	100.0	excluded; properties used instead
Sample	<code>fuel_density_kg_m3</code>	kg/m ³	100.0	predictor
Sample	<code>fuel_k_W_mK</code>	W/(m K)	100.0	predictor
Sample	<code>fuel_cp_J_kgK</code>	J/(kg K)	100.0	predictor
Sample	<code>fuel_pyrolysis_T_K</code>	K	100.0	predictor

Table 3.1: Complete field dictionary of the Excel database (continued).

Group	Excel field	Units	Compl. (%)	Modeling role
Sample	fuel_alpha_m2_s	m ² /s	100.0	derived predictor
Sample	fuel_volumetric_heat_capacity_J_m3K	J/(m ³ K)	100.0	derived predictor
Sample	core_density_kg_m3	kg/m ³	100.0	predictor
Sample	core_k_W_mK	W/(m K)	100.0	predictor
Sample	core_cp_J_kgK	J/(kg K)	100.0	predictor
Sample	core_volumetric_heat_capacity_J_m3K	J/(m ³ K)	100.0	derived predictor
Flow	Oxygen Concentration	mole fraction	100.0	predictor
Flow	diluent	categorical	100.0	predictor
Flow	gas_M	kg/mol	100.0	predictor
Flow	gas_cp_mass	J/(kg K)	100.0	predictor
Flow	gas_k	W/(m K)	100.0	predictor
Flow	gas_density_kg_m3	kg/m ³	98.9	predictor
Flow	gas_alpha_m2_s	m ² /s	98.9	derived predictor
Flow	gas_nu_m2_s	m ² /s	98.9	derived predictor
Flow	Reynolds	—	98.9	derived predictor
Flow	Peclet	—	98.9	derived predictor
Flow	Prandtl	—	98.9	derived predictor
Flow	thermal_diffusion_time	s	100.0	derived predictor
Flow	Pressure (Pa)	Pa	98.9	predictor
Flow	Flow Velocity (m/s)	m/s	79.5	predictor
Flow	Rig Name	free text	76.6	excluded apparatus
Flow	Internal Geometry	free text	86.9	excluded apparatus
Flow	Internal Dimensions	m	90.3	excluded apparatus
Flow	Gravity	g/g_{earth}	98.9	predictor
Extra	Facility	categorical	98.9	excluded apparatus
Extra	Ignition Method	categorical	98.6	predictor
Extra	Ignition Power (W)	W	18.8	predictor (sparse)
Extra	Ignition Time (s)	s	45.6	predictor (sparse)
Outputs	Ignition	Yes/No	100.0	classification target
Outputs	Flame Length	m	4.3	post-outcome; excluded
Outputs	FSR (m/s)	m/s	51.5	regression target
Outputs	HRR	W	1.1	post-outcome; excluded
Outputs	Smoke/Aerosols	—	1.8	post-outcome; excluded
Info	Info	free text	25.5	notes; excluded

Ignition labels, fuel thermophysical properties, oxygen, and diluent are essentially complete. Geometry, gravity, pressure, and ignition method remain above 98%. Flow velocity drops to 79.5%. Numeric FSR is present in only 51.5% of the raw rows, which is why the regression cohort is smaller than the ignition cohort. Flame length (4.3%), heat-release rate (1.1%), and smoke/aerosols (1.8%) are too sparse to support a third supervised task. Ignition power (18.8%) and ignition time (45.6%) are retained as predictors where reported, with fold-local imputation for the remainder. Apparatus names, facility labels, free-text material names, and notes are excluded from model inputs because they can fingerprint a laboratory without encoding physics. The schema therefore justifies both the two-task scope and the need for fold-local imputation: missingness is a property of what laboratories reported, not a random mask.

Quality control continued in code. Encodings were detected, empty rows and columns removed, target labels normalized, exact modeling duplicates removed, and unparseable values recorded as missing rather than guessed. Physical-domain checks rejected impossible concentrations, non-positive pressures, and invalid target encodings; conversions were performed before duplicate detection so that equivalent observations expressed in different units could not survive as independent rows. Canonical paper identity preferred a normalized DOI and otherwise used a normalized citation. This produces identifiers of the form `doi::<normalized-doi>` or `citation::<normalized-citation>`. Cosmetic citation differences therefore cannot place one publication in two paper groups. Dataset shape, column roles, target prevalence, paper counts, and a content fingerprint are written to validation records and checked again when frozen splits are loaded.


The two tasks use different valid subsets. After feature/target standardization and deduplication, ignition classification contains 4,487 rows from 87 canonical papers (3,383 ignition and 1,104 no ignition). FSR regression contains 2,531 rows with a numeric FSR from 66 papers. The FSR papers contribute between 1 and 193 rows, with median 22.5. Dataset fingerprints are persisted with the splits and evaluations so that a result cannot silently refer to a changed database.

3.1.4 Source-size and class imbalance

The assembled literature table is not a balanced factorial design. Two imbalances dominate every subsequent modeling choice.

The first is source-size imbalance. Publications contribute very different numbers of nearby conditions from a single rig, material batch, or campaign. Figure 3.1 already shows that the median paper is modest while a handful of campaigns dominate the row count; the Pareto curve in figure 3.2(a) makes the cumulative effect explicit. After sorting papers from largest to smallest contribution, 36 of 87 papers account for more than 80% of all rows. The remaining long tail contributes the last 20%. A loss that treats every row as an independent draw therefore over-represents densely sampled campaigns and under-represents sparse but physically distinct sources. Paper weighting, paper-grouped evaluation, and paper-cluster bootstrap intervals are the methodological response to this structure.

The second is class imbalance on the ignition label. In the raw assembled table of 5,057 rows, 3,875 observations (76.6%) are ignition and 1,182 (23.4%) are no ignition (figure 3.2(b)). After task-specific cleaning the modeling cohort remains majority-positive: 3,383 ignition versus 1,104 no ignition. A classifier that always predicts ignition would therefore already achieve about three-quarters accuracy and a precision–recall baseline equal to the positive prevalence. Accuracy is not a defensible selection metric. The ignition pipeline instead compares models with ROC–AUC, PR–AUC, and Brier score, and it reports thresholded sensitivity, specificity, precision, F1, balanced accuracy, and Matthews correlation. Class weights and paper weights are tested as candidate recipes rather than assumed to correct the imbalance automatically.



(a) Cumulative data share after sorting papers by contribution. Thirty-six papers supply more than 80% of the rows.

(b) Raw ignition outcome composition: 76.6% ignition versus 23.4% no ignition.

Figure 3.2: Two sampling imbalances in the assembled database. Panel (a) shows that a minority of publications dominate the row count; panel (b) shows that ignition is the majority class. Both motivate grouped evaluation and prevalence-aware classification metrics.

3.2 Methods common to both machine-learning pipelines

The two pipelines share a scientific contract even though one classifies ignition and the other regresses flame-spread rate. Both estimate a mapping from measured physical inputs $\mathbf{x}$ to an observed target y : the classifier estimates $\hat{p} = P(y_{\text{ign}} = 1 \mid \mathbf{x})$; the regressor estimates $\hat{v}_{\text{FSR}}$. Both freeze splits before any model is compared, nest hyperparameter search inside outer training partitions, and report interpolation and extrapolation as different questions. Table 3.2 states the shared design contract; later subsections give task-specific implementations.

Table 3.2: Pre-specified design contract for both prediction tasks. The interpolation and extrapolation columns define different target populations and are therefore reported separately.

Design element	Interpolation	Extrapolation
Target use	New condition within represented campaigns	New paper, rig, or experimental campaign
Outer assignment	Individual rows	Complete canonical papers
Paper overlap	Permitted by definition	Prohibited in inner and outer folds
Resampling unit for uncertainty	Observation after repeated-prediction aggregation	Canonical-paper cluster
Data-adaptive operations	Training-fold imputation, scaling, encoding, weighting, tuning, and threshold selection only	
Primary evidence	Persisted outer-test predictions; full-data refits are inference artifacts, not evaluation evidence	

The repeated splits estimate performance for the empirical mixture of conditions and papers represented by

the curated corpus. They do not constitute an external prospective validation set, and paper-grouped performance should not be interpreted as guaranteed transfer beyond the measured feature envelope. The subsections below explain why the model families were chosen, how nested evaluation is enforced, and which preprocessing and weighting choices are common to both tasks.

3.2.1 Benchmarking motivation and model families

The closest prior machine-learning study of flame spread is Rivera, Jose, and colleagues, who compiled approximately 1,200 electrical-wire observations and compared linear regression, elastic net, stochastic gradient descent, Bayesian ridge, kernel ridge, decision trees, k -nearest neighbours, gradient boosting, support-vector regression, and a multilayer perceptron under random ten-fold cross-validation [9]. That paper is the explicit benchmark for family choice in this report. Nonlinear models were competitive there; linear and shrinkage families lagged. The present work therefore retains the families that were scientifically informative in that study, upgrades gradient boosting to XGBoost, and adds random forest as a bagging control for FSR. Linear, elastic-net, SGD, Bayesian-ridge, and kernel-ridge regressors are not re-run: Jose already showed that they are a weak hypothesis for heterogeneous flame-spread tables, and repeating them would inflate the portfolio without a new physical question.

XGBoost is the primary tabular learner because Shwartz-Ziv and Armon found gradient-boosted trees to be a strong, readily tuned baseline on heterogeneous tables, including settings where deep models did not dominate [1, 11]. That prior is a starting point, not a result. Decision trees remain as a transparent nonlinear baseline from Jose. k -NN tests whether nearby experimental conditions are locally predictive. SVM/SVR tests linear versus radial kernels in the engineered physics space. MLPs test whether a smooth universal approximator helps once the same nested, leakage-controlled protocol is imposed. Random forest is the one family Jose did not include; it isolates bagging from boosting on the FSR task, where source shift and outliers are acute.

Rivera/Jose family	Present benchmark	Scientific role	Task
Gradient boosting	XGBoost	Strong tabular baseline; boosting	Both
Decision tree	Decision tree	Interpretable nonlinear baseline	Ignition
—	Random forest	Bagging control added in this work	FSR
k -nearest neighbours	k-NN	Local similarity in condition space	Both
Multilayer perceptron	MLP	Smooth nonlinear approximation	Both
Support-vector regression	SVM / SVR	Linear versus radial kernels	Both
Linear and shrinkage models	Not retained	Weak prior on heterogeneous FSR table	

Protocol extension: nested selection; paper-disjoint testing; source/class weighting; robust losses and target transforms.

Figure 3.3: Mapping from Rivera/Jose model families to the portfolios used here. Nonlinear families are retained or strengthened; linear families are dropped because they lagged in that wire-FSR study. Random forest is added only for FSR. Protocol differences—nested interpolation versus paper-grouped extrapolation, paper and class weights, focal loss, signed-log targets, and real-data versus bootstrap conditions—are deliberate extensions, not a claim that the two papers are numerically interchangeable [9, 11].

3.2.2 Features, leakage exclusions, and preprocessing

Paper identifiers, authors, DOI, facility names, free text, and post-outcome quantities are never used as predictors. Flame length, HRR, and smoke/aerosols are excluded from both feature sets. Material names, rig names, and internal apparatus geometry are likewise excluded: they can identify a laboratory without encoding a physical mechanism. The appendix records the deployed manifests. Ignition uses a fixed list of 32 numeric and five categorical physics features. FSR detects 29 numeric and four categorical fields from the same schema. Both cover fuel and core properties, geometry, oxygen and diluent, pressure, gas transport, flow, gravity, and ignition conditions where available. The persisted FSR manifest also contains categorical **Ignition**; because this conflicts with explanatory documentation, it is disclosed as a schema limitation rather than hidden.

Numeric missing values are replaced by the training-fold median for models that cannot handle missingness; XGBoost retains native numeric missingness. Distance, margin, and neural models use training-fold standardiza-

tion. Categorical missingness becomes an explicit “Unknown” level, and one-hot encoding ignores a category first seen in a test fold. No vocabulary, scaler, imputer, weight, or threshold is fit on outer test rows.

3.2.3 Nested evaluation

For a reader new to machine learning, three data roles must remain distinct:

Training data fit model parameters, such as tree splits or neural weights.

Validation data choose hyperparameters, such as tree depth, regularization, or neighbourhood size.

Test data measure the selected procedure once on observations not used for either fitting or selection.

The project implements this separation with nested evaluation (figure 3.4). For outer split s , candidate c and hyperparameter setting λ , the inner loop selects

$$\lambda_{c,s}^* = \arg \min_{\lambda \in \Lambda_c} \frac{1}{K_{\text{in}}} \sum_{k=1}^{K_{\text{in}}} L\left(y^{(k)}, \hat{f}_{c,\lambda}^{(-k)}(\mathbf{x}^{(k)})\right),$$

with the direction of L reversed for metrics that are maximized. The selected configuration is refitted on the complete outer training partition and evaluated once on the untouched outer test partition. Preprocessing is part of $\hat{f}$, not a preliminary operation on the full table. Candidate comparison is paired on identical outer assignments, reducing avoidable split-to-split noise. Repeating the outer split quantifies sensitivity to data partitioning; repeated-fold scores are correlated and are not treated as independent experimental replications.




Figure 3.4: Nested evaluation contract shared by both pipelines. Inner cross-validation and threshold or recipe selection never see the outer test partition. The outer test is scored once per fold; those scores are the comparison evidence.

3.2.4 Interpolation versus extrapolation

Both pipelines answer two different questions, illustrated in figure 3.5:

- **Interpolation** holds out rows while allowing other rows from the same paper to remain in training. It asks whether the model predicts new conditions inside campaigns it partly knows.
- **Extrapolation** holds out complete canonical papers. It asks whether the procedure transfers to a new publication, rig, or campaign and is the primary generalization test.

The distinction prevents paper-specific sampling and apparatus signatures from being mistaken for broad physical transfer. Rivera et al. themselves describe random cross-validation as robustness within a compiled table rather than strict held-source transfer [9]. The present design makes that caveat operational: all candidates within a task share the same frozen outer assignments, and grouped protocols assert zero paper overlap.

3.2.5 Paper weighting, uncertainty, and refitting

Source-size imbalance is treated as a modeling hypothesis, not a background correction that is assumed to work. For loss ℓ , paper $p(i)$ associated with row i , and training-fold paper size n_p , the weighted empirical risk is

$$\hat{R}_w(f) = \frac{\sum_{i \in \mathcal{T}} w_{p(i)} \ell(y_i, f(\mathbf{x}_i))}{\sum_{i \in \mathcal{T}} w_{p(i)}}.$$




Figure 3.5: Schematic of the two evaluation questions. Interpolation mixes train and test rows inside the same papers. Extrapolation holds out entire papers, which is the relevant test for a new campaign or material family.

Both pipelines test three strategies (figure 3.6). Unweighted training, $w_p = 1$, gives paper p total influence proportional to n_p . Square-root weighting, $w_p = n_p^{-1/2}$, reduces that influence to $\sqrt{n_p}$. Inverse weighting, $w_p = n_p^{-1}$, gives each paper approximately equal total influence. Counts and normalized weights are always computed from the current training partition. Families that cannot accept sample weights are trained on a deterministic probability-proportional resample of that partition so that a weighting flag is never silently ignored.




Figure 3.6: Paper-weighting strategies shared by ignition and FSR candidates. n_p is the number of training rows from paper p . The preferred strategy is selected empirically from nested scores, not assumed a priori.

Outer-fold means and standard deviations summarize partition sensitivity. Ninety-five-percent percentile intervals use 2,000 bootstrap replicates. Repeated outer predictions are first averaged for each observation; the bootstrap then resamples observations for interpolation and complete canonical paper clusters for extrapolation. This avoids counting repeated predictions of the same row as independent evidence and preserves the sampling unit of the question being asked. Intervals quantify resampling uncertainty within the assembled corpus; they do not include literature-selection, transcription, or measurement uncertainty.

Only after selection is complete is the chosen configuration refitted on all labeled rows to create an inference artifact. That full-data model is deployable, but its training fit is not used as evidence of generalization. Champion claims apply only to candidates present in the persisted selection evidence: the merged FSR comparison contains 25 of the 34 configured candidates, and the formal ignition selection record omits available MLP shards. This distinction prevents an unevaluated or incomplete candidate from becoming a champion by configuration alone.

3.3 Ignition-classification pipeline

3.3.1 Task, probabilities, and metrics

Ignition is a binary label, but the model returns a probability $\hat{p} \in [0, 1]$. A threshold τ produces the decision

$$\hat{y}(\tau) = \begin{cases} 1, & \hat{p} \geq \tau, \\ 0, & \hat{p} < \tau. \end{cases}$$

A lower threshold catches more ignitions and usually increases false alarms; a higher threshold does the reverse. Because the assembled table is majority-positive, accuracy is not a defensible selection metric. Comparison therefore uses ranking and calibration scores that do not require a threshold: ROC-AUC, PR-AUC, and the Brier score

$$\text{BS} = n^{-1} \sum_{i=1}^n (\hat{p}_i - y_i)^2.$$

ROC-AUC measures pairwise ranking across classes; PR-AUC is interpreted relative to the positive-class prevalence; and Brier score assesses squared probability error. Four operating points are then frozen from inner out-of-fold predictions—maximum Matthews correlation (MCC), maximum F1, maximum balanced accuracy, and maximum Youden $J = \text{sensitivity} + \text{specificity} - 1$ —and applied to the outer test fold without re-optimizing on test labels. Thresholded reporting adds sensitivity, specificity, precision, F1, balanced accuracy, and MCC. The threshold objective is therefore part of the procedure being evaluated, not a post-hoc choice made after observing the test set.

3.3.2 Candidate design

The ignition portfolio contains 26 candidates across five families (figure 3.7). The design is an ablation of *how* a family is trained, not a search over dozens of unrelated architectures. Every family includes an unweighted baseline. XGBoost, as the primary model, receives the full imbalance grid; the other families are structured Jose-aligned baselines.

Two imbalance mechanisms are treated as orthogonal. Paper weighting corrects sampling bias from large campaigns. Class weighting rescales the loss by label frequency. Focal loss [5] down-weights easy, confident examples dynamically and is therefore an alternative class-imbalance mechanism, not a third knob to stack on class weights. Focal candidates use $\gamma = 2$ and `class_weight: false` to avoid double-correcting the majority class. Paper weighting is retained with focal loss because it corrects a different bias.




Figure 3.7: Ignition candidate portfolio. XGBoost carries the full paper-weight $\times$ class-weight ablation plus standalone focal loss. Decision tree, k -NN, MLP, and SVM supply comparison baselines with one structural variant each.

The remaining 18 candidates test one structural hypothesis per family (table 3.4). Decision-tree depth, leaf size, and split criterion are already inner-search hyperparameters, so no extra “pruned tree” candidate is added. Distance-weighted k -NN lets closer experimental neighbours vote more strongly. Linear SVM tests whether the ignition boundary is linearly separable in the physics feature space, which is plausible if oxygen and pressure thresholds dominate. MLP focal variants mirror the XGBoost focal logic at reduced scale.

Table 3.3: XGBoost ignition ablation (8 candidates). The first six cells are the 2×3 grid of paper weight and class weight under logistic loss. The last two isolate focal loss.

Candidate	Paper weight	Class weight	Objective
xgb_baseline	none	no	logistic
xgb_paper_sqrt	sqrt	no	logistic
xgb_paper_inverse	inverse	no	logistic
xgb_class_only	none	yes	logistic
xgb_weighted_sqrt	sqrt	yes	logistic
xgb_weighted_inverse	inverse	yes	logistic
xgb_focal_sqrt	sqrt	no	focal $\gamma = 2$
xgb_focal_inverse	inverse	no	focal $\gamma = 2$

Table 3.4: Non-XGBoost ignition candidates (18). Combined weighting is used for comparison families; structural variants are crossed with both paper-weight strategies.

Candidate	Family	Paper / class weight	Structural variant
dt_baseline	decision tree	none / no	—
dt_weighted_sqrt	decision tree	sqrt / yes	—
dt_weighted_inverse	decision tree	inverse / yes	—
knn_baseline	k -NN	none / no	uniform neighbours
knn_weighted_sqrt	k -NN	sqrt / yes	uniform neighbours
knn_weighted_inverse	k -NN	inverse / yes	uniform neighbours
knn_distance_sqrt	k -NN	sqrt / yes	distance weights
knn_distance_inverse	k -NN	inverse / yes	distance weights
mlp_baseline	MLP	none / no	—
mlp_weighted_sqrt	MLP	sqrt / yes	—
mlp_weighted_inverse	MLP	inverse / yes	—
mlp_focal_sqrt	MLP	sqrt / no	focal $\gamma = 2$
mlp_focal_inverse	MLP	inverse / no	focal $\gamma = 2$
svm_baseline	SVM	none / no	RBF kernel
svm_weighted_sqrt	SVM	sqrt / yes	RBF kernel
svm_weighted_inverse	SVM	inverse / yes	RBF kernel
svm_linear_sqrt	SVM	sqrt / yes	linear kernel
svm_linear_inverse	SVM	inverse / yes	linear kernel

Hyperparameters that are *not* part of a candidate identity are sampled inside each outer training partition (table 3.5). Inner ranking uses mean ROC–AUC, then PR–AUC, then lower Brier score. This keeps the candidate list interpretable—each ID is a scientific hypothesis—while still allowing each family a fair computational budget.

Table 3.5: Inner-search spaces for ignition families. Candidate YAML fixes only the hypothesis under test (weights, kernel, focal objective); the inner loop samples these knobs.

Family	Sampled hyperparameters
XGBoost	$n_{\text{estimators}} \in \{200, 400, 600, 800\}$; learning rate $\{0.01, 0.03, 0.05, 0.1\}$; depth $\{2, \dots, 6\}$; min-child-weight, subsample, column sample, ℓ_1/ℓ_2 regularisation, gamma.
Decision tree	depth $\{3, 5, 7, 10, \infty\}$; min-leaf $\{1, 2, 5, 10, 20\}$; split criterion gini, entropy, or log-loss.
k -NN	$k \in \{5, 9, 15, 25, 40, 60\}$; uniform or distance weights unless fixed; Minkowski $p \in \{1, 2\}$; leaf size.
MLP	hidden widths (32), (64), (128, 64), (128, 64, 32); ℓ_2 and learning-rate grid; ReLU; batch size 32 or 128.
SVM	$C \in \{0.1, 0.3, 1, 3, 10, 30\}$; RBF γ grid or linear kernel when fixed; shrinking on/off.

3.3.3 Pipeline architecture

Figure 3.8 is the executable workflow. After the feature manifest is built, four split protocols are frozen once and reused by every candidate: stratified interpolation cross-validation, a stratified 80/20 holdout (consistency check only), paper-grouped extrapolation, and leave-one-paper-out (LOPO). For each candidate and outer fold the inner search runs on training rows only, four thresholds are frozen from inner out-of-fold probabilities, and the outer test fold is scored once. Selection reads persisted tables and does not refit. Only the winning interpolation and extrapolation configurations are then trained on all labeled rows for later inference.

Ignition classification — executable evaluation design




Figure 3.8: Ignition classification pipeline. Splits are frozen before model comparison. Nested search and threshold selection are confined to outer training partitions. LOPO is a mandatory robustness gate for the extrapolation champion, not a substitute for grouped nested CV.

3.3.4 Interpolation protocol

Ignition interpolation uses five-fold stratified cross-validation repeated for seeds 42–44, giving 15 outer folds. Stratification keeps the ignition/no- ignition proportion similar in each fold while rows are assigned independently, so papers may appear on both sides. Within each outer training partition, stratified inner folds search the family space. A separate stratified 80/20 holdout for each seed is a quick consistency check, not the primary selection protocol.

Interpolation selection requires valid probabilities, complete folds, passed integrity checks, and ROC–AUC fold standard deviation no greater than 0.15. Candidates within 0.01 of the best ROC–AUC are resolved by higher PR–AUC, lower uncertainty, and a declared simplicity order (decision tree, k -NN, SVM, XGBoost, MLP). The resulting champion is refitted on all ignition-labeled rows for within-domain inference.

3.3.5 Extrapolation and leave-one-paper-out

Ignition extrapolation uses `StratifiedGroupKFold`: all rows from one paper stay together while each fold retains both classes where feasible. Five paper-grouped folds repeated across three seeds again produce 15 outer tests. The inner search is also paper-disjoint, so a paper held out for inner validation cannot appear in the corresponding inner training fold.

The primary metric remains ROC–AUC, but eligibility additionally requires a successful LOPO analysis. Each of the 87 papers is held out once. A single-class paper cannot define its own ROC–AUC, so LOPO is interpreted through pooled predictions, calibration, and the subset of papers for which a metric is defined. The selected extrapolation model and its thresholds are refitted on the full labeled cohort only after this paper-disjoint comparison is frozen.

3.4 Flame-spread-rate regression pipeline

3.4.1 Task and metrics

Regression predicts a continuous FSR in m/s. There is no classification threshold: quality is measured directly in target units. For outer-test residuals $e_i = \hat{y}_i - y_i$,

$$\text{RMSE} = \sqrt{n^{-1} \sum_i e_i^2}, \quad \text{MAE} = n^{-1} \sum_i |e_i|.$$

RMSE penalizes larger errors more strongly and is the primary selection metric. MAE reports typical absolute error, R^2 compares residual variance with a mean predictor, and mean bias error $n^{-1} \sum_i e_i$ records systematic over- or underprediction. MAPE is retained only as a secondary diagnostic because exact and near-zero FSR values make percentage errors unstable. Predictions are always inverted back to original FSR units before scoring, including for signed-log models.

Jose’s headline $R^2 > 0.9$ values were obtained on a narrower wire table under random ten-fold cross-validation [9]. The present FSR pipeline is built so that an analogous interpolation number can be reported, but the scientific claim of this project is the paper-grouped result. The two must not be averaged or substituted for one another.

3.4.2 Candidate design

The FSR portfolio contains 34 candidates across five families (figure 3.9). Each candidate tests one hypothesis: does paper weighting improve transfer; does a robust loss help with experimental outliers; does compressing a skewed target help optimization; is local similarity enough; is the physics feature space already linear? Jose’s linear families are omitted. Random forest is added as the bagging counterpart to XGBoost so that any later preference for trees is not confounded with boosting.

XGBoost (10 candidates) is again the most granular family. Three objectives are compared under square-root and inverse paper weights: squared error, absolute error, and pseudo-Huber. Absolute error reduces sensitivity to large residuals. Pseudo-Huber is smooth near zero and linear in the tails. Two further candidates predict $\text{sign}(y) \log(1 + |y|)$ and invert the transform at inference, testing whether compressing the heavy FSR tail improves optimization. One paper-bagging candidate averages ten XGBoost models trained on paper-level bootstrap samples, reducing dependence on any single campaign.

Random forest (7) mirrors the scientifically interesting XGBoost knobs without boosting: baseline, sqrt/inverse paper weights, absolute-error split criterion, and signed-log targets. k -NN (6) crosses uniform versus distance neighbours with none/sqrt/inverse paper weights. MLP (5) includes baseline and paper-weighted models plus two signed-log networks, because neural optimization is often more stable on compressed targets when the response spans orders of magnitude. SVR (6) crosses linear versus RBF kernels with the same three paper weights: a strong linear kernel would imply that feature engineering already captured most of the spread physics.

Inner search again samples family hyperparameters rather than candidate identity (table 3.7). Configurations are ranked by mean inner RMSE, then MAE, then higher R^2 .

3.4.3 Pipeline architecture

Figure 3.10 shows the regression workflow. After role detection and leakage exclusions, interpolation and extrapolation holdouts are frozen. Each candidate is evaluated as an independent shard: nested search on the outer training partition, then two fits of the selected configuration. The `rd` condition uses real training rows only. The




Figure 3.9: FSR candidate portfolio. XGBoost explores loss, target scale, paper weights, and paper bagging. Random forest is the bagging control. *k*-NN, MLP, and SVR supply local, neural, and kernel baselines aligned with Jose.

Table 3.6: Configured FSR candidates (34). **rd** and **rd_bt** conditions are applied to every candidate at evaluation time, so the comparison pool is a pair (candidate, augmentation).

Candidate	Family	Paper weight	Hypothesis
<code>xgb_baseline</code>	XGBoost	none	squared-error baseline
<code>xgb_paper_sqrt / _inverse</code>	XGBoost	<code>sqrt / inverse</code>	paper reweighting
<code>xgb_absoluteerror_*</code>	XGBoost	<code>sqrt / inverse</code>	robust ℓ_1 objective
<code>xgb_pseudohuber_*</code>	XGBoost	<code>sqrt / inverse</code>	smooth robust objective
<code>xgb_log_target_*</code>	XGBoost	<code>sqrt / inverse</code>	signed-log($1 + y $) target
<code>xgb_paper_bagging_sqrt</code>	XGBoost	<code>sqrt</code>	paper-level bagging ($B = 10$)
<code>rf_baseline</code>	random forest	none	bagged-tree baseline
<code>rf_paper_sqrt / _inverse</code>	random forest	<code>sqrt / inverse</code>	paper reweighting
<code>rf_absoluteerror_*</code>	random forest	<code>sqrt / inverse</code>	MAE split criterion
<code>rf_log_target_*</code>	random forest	<code>sqrt / inverse</code>	signed-log target
<code>knn_uniform*</code>	<i>k</i> -NN	<code>none / sqrt / inverse</code>	equal neighbour votes
<code>knn_distance*</code>	<i>k</i> -NN	<code>none / sqrt / inverse</code>	distance-weighted votes
<code>mlp_baseline / paper_*</code>	MLP	<code>none / sqrt / inverse</code>	neural baseline
<code>mlp_log_target_*</code>	MLP	<code>sqrt / inverse</code>	signed-log target
<code>svr_linear*</code>	SVR	<code>none / sqrt / inverse</code>	linear kernel
<code>svr_rbf*</code>	SVR	<code>none / sqrt / inverse</code>	RBF kernel

Table 3.7: Inner-search spaces for FSR families.

Family	Sampled hyperparameters
XGBoost	$n_{\text{estimators}} \in \{100, 200, 300, 500, 800\}$; learning rate $\{0.01, 0.02, 0.05, 0.1, 0.2\}$; depth $\{2, \dots, 6\}$; min-child-weight, subsample, column sample, ℓ_1/ℓ_2 , gamma.
Random forest	100–500 trees; depth 4–12 or unlimited; split and leaf minima; max_features in $\{\sqrt{\cdot}, \log_2, 0.5, 0.8, 1\}$.
<i>k</i> -NN	$k \in \{3, 5, 7, 9, 11, 15, 21, 31\}$; uniform or distance unless fixed; $p \in \{1, 2\}$.
MLP	widths including (100), (128, 64), (256, 128); ReLU or tanh; ℓ_2 and learning-rate grid; batch size 32–128.
SVR	$C \in \{0.1, \dots, 100\}$; $\varepsilon \in \{0.01, \dots, 0.5\}$; RBF γ grid or linear kernel when fixed.

`rd_bt` condition appends 1,000 training rows sampled with replacement; copies inherit their source paper identity and never enter the test partition. Bootstrap augmentation is therefore a reweighting of observed conditions, as in Jose’s most conservative augmentation finding, not a generator of new experiments [9]. Shards are merged before selection so that champions are chosen from tables, not from a second round of training.

Flame-spread-rate regression — executable evaluation design




Figure 3.10: FSR regression pipeline. Every candidate is scored under real-data (`rd`) and bootstrap-augmented (`rd_bt`) training. Test partitions are never augmented. Selection operates on merged, persisted metrics.

3.4.4 Interpolation protocol

FSR interpolation uses ten repeated, approximately 80/20 row holdouts with seeds 42–51. Target values are binned into six quantiles where feasible so that each split represents the strongly skewed FSR distribution. Papers may occur on both sides. Inside each outer training partition, shuffled K -fold validation compares hyperparameters. The selected recipe is fitted on all 2,024 training rows and tested on the 507 untouched rows for that repeat.

The selection policy rejects incomplete or non-finite evaluations, RMSE coefficient of variation above 0.45, and mean $R^2 < -0.25$. Candidates within 2% of the lowest RMSE form a tie band resolved by higher R^2 , lower uncertainty, preference for real data over bootstrap copies, simpler family (random forest, k -NN, SVR, XGBoost, MLP), and fewer features where the feature sets differ.

3.4.5 Extrapolation protocol

FSR extrapolation also uses ten repeated 80/20 holdouts, but the unit assigned to train or test is the paper. For each repeat, 50 candidate `GroupShuffleSplit` partitions are generated and the most target-balanced paper-disjoint split is retained. Each test contains 14 papers and each training partition 52; row counts vary with paper size. Balance screening reduces avoidable target-range changes without relaxing paper separation. Because the screening criterion uses only target-distribution summaries to choose among pre-generated paper partitions, it is treated as part of split construction and is frozen before any candidate is evaluated. The result remains a transfer estimate for this balance-screened split policy, not for an arbitrary future collection of papers.

Inner tuning uses `GroupKFold` by paper. Extrapolation tie-breaks add the interpolation–extrapolation gap so that a recipe is not favoured solely for excellent within-paper fit. Cluster bootstrap intervals resample papers rather than thousands of correlated rows. The winning (candidate, augmentation) pair is refitted on all 2,531 FSR rows and stored separately from the outer-fold evidence.

3.5 Reproducibility and audit trail

Reproducibility is enforced through persisted interfaces rather than relying on narrative description alone. The cleaned task table is fingerprinted; outer and inner assignments are generated once from declared seeds and stored with canonical paper identifiers; candidate hypotheses and search spaces are versioned configuration files;

and each evaluation shard records software settings, selected hyperparameters, fold metrics, and row-level outer predictions. Aggregation and champion selection read these artifacts without training a model. Integrity checks require complete expected folds, finite metrics, valid probability ranges for classification, and zero paper overlap in grouped protocols. The final refit stores its feature manifest, preprocessing objects, candidate identity, selection provenance, and data fingerprint beside the serialized estimator. Consequently, a reported number can be traced back to an immutable prediction ledger, while an inference artifact can be rejected if its schema or database fingerprint does not match.

The figure source is similarly reproducible. Figures 3.3–3.10 are generated by `Report/figures/generate_method_charts.py` as vector PDFs with embedded fonts; high-resolution PNG files are retained only as review fallbacks. This avoids resolution loss and ensures consistent typography, colour, line weight, and physical dimensions in the compiled report.

3.6 Explainability and physical interpretation

Explainability is performed after evaluation; it does not influence model selection. Tree impurity importance summarizes how often and how effectively a feature splits training data. Permutation importance measures the loss of held-out performance after one original feature is shuffled. SHAP decomposes each prediction into additive feature contributions relative to a background:

$$f(\mathbf{x}_i) = \phi_0 + \sum_j \phi_{ij}, \quad I_j^{\text{SHAP}} = \frac{1}{n} \sum_i |\phi_{ij}|.$$

Beeswarm plots retain both sign and feature value, whereas mean absolute SHAP bars show magnitude only.

These are explanations of fitted functions, not causal experiments. Correlated raw and derived features can substitute for one another; missingness can encode paper identity; and a physically plausible ranking does not prove a governing mechanism. Importance is therefore reported with the model, task, and evaluation regime that produced it rather than as a universal material ranking. The report compares SHAP with held-out permutation importance and interprets concordant patterns as hypotheses for testing and data collection, not as certification evidence. No feature is described as a causal driver solely because it receives a large attribution.

4 Results

After the preprocessing described in chapter 3, the ignition cohort contains 4,487 labeled rows from 87 papers and the FSR cohort contains 2,531 numeric observations from 66 papers. The figures below report those modeling tasks. Database sampling geometry that is not itself a predictive score—in particular oxygen–pressure coverage—is interpreted in section 5.3 as a map of the experimental search space that remains to be filled.

4.1 Ignition classification

4.1.1 ROC, precision–recall, and the fire-safe operating question

A binary ignition model does not output a physical yes/no switch. It outputs a score $\hat{p} \in [0, 1]$, interpreted as an estimated probability of ignition. Converting that score into a decision requires a threshold τ :

$$\hat{y}(\tau) = \begin{cases} 1, & \hat{p} \geq \tau, \\ 0, & \hat{p} < \tau. \end{cases} \quad (4.1)$$

Each τ produces a confusion-matrix count of true positives (TP), false positives (FP), true negatives (TN), and false negatives (FN). The elementary rates used throughout this section are

$$\text{TPR} = \text{sensitivity} = \text{recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}, \quad (4.2)$$

$$\text{FPR} = 1 - \text{specificity} = \frac{\text{FP}}{\text{FP} + \text{TN}}, \quad (4.3)$$

$$\text{precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}. \quad (4.4)$$

Recall is the fraction of true ignitions that are caught. Precision is the fraction of predicted ignitions that actually ignite. Specificity is the fraction of true non-ignitions that are left unlabeled. No single rate is a complete account of a fire-safety call: missing an ignition (false negative) is the safety-critical error, while flooding a test queue with false alarms (false positives) makes the screener unusable.

Two threshold-independent curves summarize the whole family of operating points. The receiver-operating-characteristic (ROC) curve plots TPR against FPR as τ varies. A no-skill classifier that ranks at random traces the diagonal, with area $1/2$. The area under the ROC curve (ROC–AUC) equals the probability that a randomly chosen ignition receives a higher score than a randomly chosen non-ignition. It is a ranking metric: it does not select τ , and it weights the two classes symmetrically through TPR and FPR.

The precision–recall (PR) curve plots precision against recall. Its no-skill baseline is not $1/2$; it is the positive prevalence $\pi = P(y = 1)$. In this database $\pi \approx 0.74$ – 0.75 , so an always-ignite rule already attains precision equal to that prevalence and recall equal to one. The area under the PR curve (PR–AUC, equivalently average precision in figure 4.1) therefore measures how much ranking quality exists *above* that naive majority policy.

A convenient single-threshold summary of the PR trade-off is the F1 score, the harmonic mean of precision and recall,

$$\text{F1} = \frac{2 \text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}. \quad (4.5)$$

Because a harmonic mean is dominated by the smaller of its two arguments, F1 is large only when both precision and recall are large. It does not, however, use true negatives, so a majority-positive task can reward a low τ that labels almost everything as ignition. Complementary summaries that do use the full confusion matrix are balanced accuracy (sensitivity + specificity)/2, Youden’s $J = \text{sensitivity} + \text{specificity} - 1$, and the Matthews correlation

$$\text{MCC} = \frac{\text{TP TN} - \text{FP FN}}{\sqrt{(\text{TP} + \text{FP})(\text{TP} + \text{FN})(\text{TN} + \text{FP})(\text{TN} + \text{FN})}}. \quad (4.6)$$

Thresholds that maximize these criteria are selected inside training folds and are interpreted as operational limitations in section 5.1; they are not used to claim a unique “best” fire-safety setting.

Figure 4.1 shows both curves for the selected XGBoost classifier on a held-out partition. The ROC curve rises steeply toward the upper-left corner and attains AUC 0.897, well above the chance diagonal. That result says the model ranks ignitions above non-ignitions for most pairs. It does *not* by itself say that a fire-safe call is available: a high ROC–AUC can coexist with many false positives when negatives are the minority class, because FPR is computed only among true non-ignitions.

The PR curve is the more informative fire-safety diagnostic for exactly that reason. Precision remains near one up to recall of about 0.7, then declines as the threshold is lowered to recover the remaining ignitions. Average

precision is 0.963, far above the prevalence baseline of 0.74 drawn as the dashed line. In operational language: there exists a high-precision region in which flagged cases are very likely to be true ignitions, and the model can recover a large majority of ignitions before precision collapses. Pushing recall toward one—the instinctive “never miss a fire” policy—still costs precision, which is why PR-AUC, rather than ROC-AUC or accuracy, is the curve that should govern a screening claim on this imbalanced label. ROC-AUC remains useful as a threshold-free ranking score for model comparison; PR-AUC is the quantity that asks whether positive-class decisions remain reliable.




Figure 4.1: Held-out ROC and precision-recall curves for the ignition XGBoost classifier. The ROC panel (left) plots true-positive rate against false-positive rate; the shaded area is ROC-AUC = 0.897, against a chance diagonal of 0.5. The PR panel (right) plots precision against recall; the shaded area is average precision = 0.963, against a no-skill baseline equal to the positive prevalence (0.74). The PR curve is the more relevant fire-safety summary because ignition is the majority class and because a screening call is a statement about the positive class.

4.1.2 Interpolation and extrapolation by model family

The selected interpolation classifier is the paper-weighted XGBoost recipe `xgb_paper_sqrt`. Across 15 repeated stratified outer folds it reaches ROC-AUC 0.9076 ± 0.0102 , PR-AUC 0.9687 ± 0.0039 , and mean Brier score 0.1063 . The selected extrapolation classifier is the unweighted `xgb_baseline`. Its paper-grouped ROC-AUC is 0.7183 ± 0.0444 , PR-AUC 0.8767 ± 0.0298 , and Brier score 0.1791 . The paper-cluster bootstrap estimate of grouped ROC-AUC is 0.7101 with 95% interval 0.6408 – 0.7848 ; leave-one-paper-out robustness is similar (0.7020 , 0.6376 – 0.7751).

Figure 4.2 presents the same two questions for the best persisted candidate in each family, so interpolation and extrapolation are read together rather than as separate leaderboards. Panel (a) uses ROC-AUC; panel (b) uses PR-AUC. In every family the interpolation marker lies to the right of the extrapolation marker, with much smaller error bars. XGBoost is the strongest family on ROC-AUC for both protocols (interpolation 0.901 ± 0.002 , extrapolation 0.718 ± 0.044 on the plotted representatives). SVM and MLP occupy a middle band; decision trees and k -nearest neighbors are weaker, and the decision-tree extrapolation bar is the most unstable ($SD = 0.100$).

PR-AUC preserves the same family ranking for interpolation, with XGBoost at 0.966 ± 0.001 . Under extrapolation the picture is slightly different: SVM reaches 0.879 ± 0.022 and XGBoost 0.877 ± 0.030 , so the two families are indistinguishable on the positive-class ranking metric once papers are held out. That is consistent with the fire-safety reading of figure 4.1: ROC-AUC and PR-AUC can disagree about a close call because they answer different questions. All families remain well above the prevalence baseline of ≈ 0.75 , but the drop from interpolation (PR-AUC ≈ 0.93 – 0.97) to extrapolation (≈ 0.82 – 0.88) is the quantity that limits transfer to a new campaign. Tree-depth flexibility (XGBoost versus a single decision tree) and margin-based models (SVM) retain more of that positive-class ranking than memory-based KNN or an unpruned tree.

Probability ranking is not calibration. Figure 4.3 compares mean predicted probability with the observed ignition fraction. The interpolation champion tracks the identity line closely, especially above probability 0.5. The extrapolation champion is miscalibrated: its scores occupy a narrower, higher band, and high predicted probabilities overstate the observed ignition rate. A numerical probability from the grouped model therefore cannot be read as a frequentist ignition risk without a source-conditional calibration step.




Figure 4.2: Ignition interpolation (circles) versus paper-grouped extrapolation (triangles), mean $\pm$ standard deviation. Every family loses performance when entire papers are withheld, and extrapolation uncertainty is systematically larger. XGBoost leads ROC-AUC in both protocols; under PR-AUC the extrapolated SVM and XGBoost candidates are essentially tied. The two panels together replace a separate “generalization-gap” ranking: the gap is the horizontal separation between the two markers of each family.




Figure 4.3: Reliability diagram for the interpolation and extrapolation ignition champions. Perfect calibration is the dashed diagonal. The interpolation curve follows the identity; the extrapolation curve is compressed toward high scores and lies below the diagonal at the upper end, indicating over-confident high-risk predictions on held-out papers.

4.1.3 Ignition explainability

The held-paper XGBoost SHAP plot in figure 4.4 shows both contribution direction and feature value. Oxygen has the largest global SHAP magnitude, followed by absolute flow speed, Prandtl number, gas molar mass, fuel pyrolysis temperature, and ignition time. Higher oxygen values generally move predictions toward ignition in this fitted model. Low absolute flow speed produces large negative SHAP values, consistent with a quenching branch in which weak transport cannot sustain a flame. Long recorded ignition times also push the score downward, while several transport quantities show nonlinear or mixed patterns.

Permutation importance gives a different ranking: fuel pyrolysis temperature, ignition time, fuel heat capacity, and thermal diffusion time produce the largest positive held-out ROC–AUC losses when shuffled. Oxygen has high SHAP magnitude but is not the leading permutation feature in this held-out sample. The disagreement is evidence of correlated and substitutable predictors, not a contradiction and not a causal ranking. Oxygen, flow, and thermal time scales nonetheless recur across both views, which is the level of interpretation the observational table can support.




Figure 4.4: Ignition explainability on held-out papers. The beeswarm retains sign and feature value; the permutation bars measure the loss of ROC–AUC after shuffling one original feature. Repeated oxygen, flow, and thermal variables are predictive associations in the fitted function, not identified mechanisms.

4.2 Flame-spread-rate regression

4.2.1 Interpolation and extrapolation together

Regression has no classification threshold: the model predicts a continuous FSR in m/s, and quality is measured in those units. Two summaries are used throughout. RMSE is the root-mean-square error

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}, \quad (4.7)$$

which penalizes large misses and remains dimensionally an FSR. The coefficient of determination

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2} \quad (4.8)$$

compares residual variance with a mean predictor; $R^2 = 1$ is a perfect fit and $R^2 = 0$ is no better than predicting $\bar{y}$. The two metrics are not interchangeable. RMSE is the engineering error in physical units; R^2 is a scale-free statement about explained variance and is sensitive to the target spread of the held-out papers. Ranking models by RMSE can therefore disagree with ranking them by R^2 when folds differ in variance, which is why both axes appear below.

The unaugmented random-forest baseline is selected for both protocols, with separately tuned configurations. Across ten repeated row holdouts it reaches $R^2 = 0.8790 \pm 0.0178$, RMSE 0.006263 ± 0.000472 m/s, and MAE 0.002103 ± 0.000173 m/s. Across ten paper-grouped holdouts the same family reaches $R^2 = 0.4576 \pm 0.2215$, RMSE 0.012205 ± 0.004115 m/s, and MAE 0.004819 ± 0.001144 m/s. Grouped RMSE is 1.95 times the interpolation value; the paper-cluster 95% interval for grouped RMSE is 0.006525 – 0.017839 m/s. Bootstrap-augmented XGBoost has

a numerically slightly lower interpolation RMSE (0.006236 m/s) but falls within the 2% tie band and has greater fold variability, so the declared tie policy favors the unaugmented random forest.

Figure 4.5 shows every persisted candidate under both protocols and both training conditions (real data versus bootstrap-augmented training). Interpolation bars (blue) are consistently lower than grouped bars (orange). Tree ensembles occupy the left-hand, low-RMSE group; KNN occupies a middle band; SVR and MLP occupy the right-hand, high-error group, with grouped error bars that can exceed 0.03 m/s. Augmentation does not close the protocol gap: the dark and light bars of a given protocol remain close relative to the blue–orange separation. The figure therefore answers interpolation and extrapolation in one comparison: family identity and evaluation protocol dominate recipe-level tweaks.




Figure 4.5: FSR outer-fold RMSE for all persisted candidates. Blue bars are row-level interpolation; orange bars are paper-grouped extrapolation; paler and darker shades distinguish real versus bootstrap-augmented training. Tree ensembles form a leading cluster; KNN, SVR, and MLP form successively weaker clusters. The interpolation–extrapolation gap is larger than the gap among leading tree recipes.

Figure 4.6 makes the two ranking metrics explicit. Each point is a candidate; the horizontal axis is mean R^2 (higher is better) and the vertical axis is mean RMSE (lower is better), so the desirable corner is the bottom right. A dashed line separates grouped extrapolation (left, purple) from random interpolation (right, blue). Three facts follow immediately.

First, the two metrics largely agree inside each protocol: points fall along a descending cloud, so improving R^2 generally reduces RMSE. There are nevertheless two legitimate ways to rank models. RMSE is the quantity to quote for a physical error budget; R^2 is the quantity to quote when asking whether the model beat a mean predictor on that fold’s target spread. A candidate can look competitive on one axis and less so on the other if its errors concentrate on high-variance papers.

Second, model families are visually separated in both panels. Random forest and XGBoost occupy the lower-right cluster of each protocol; KNN occupies an intermediate cluster; SVR and MLP occupy the upper-left cluster. Family identity is therefore a stronger organizing principle than individual hyperparameter variants. The highlighted random-forest baseline ($R^2 = 0.879$, RMSE = 0.00626 m/s in interpolation; $R^2 = 0.458$, RMSE = 0.01221 m/s in extrapolation) sits at the front of the tree cluster in both tasks.

Third, the protocol shift is a translation of the entire family structure, not a reordering. The same relative ranking survives paper hold-out, but the absolute location moves from $R^2 \approx 0.4\text{--}0.9$ to $R^2 \approx 0.05\text{--}0.46$ and from RMSE $\approx 0.006\text{--}0.015$ to RMSE $\approx 0.012\text{--}0.016$ m/s. Interpolation competence does not imply transfer; family ranking, however, is stable enough that a screening tool can prefer tree ensembles without pretending that within-paper $R^2 = 0.88$ is available for a new campaign.

4.2.2 Held-out parity and error structure

The two panels of figure 4.7 diagnose where the grouped random forest succeeds and where it fails. Panel (a) is the extrapolation parity plot: predicted versus experimental FSR, with the one-to-one line as the reference. The pooled snapshot on that figure is $R^2 = 0.425$ and RMSE = 0.0129 m/s, consistent with the repeated-holdout mean $R^2 = 0.458 \pm 0.222$. Dense low-FSR observations lie near the diagonal. Above about 0.05 m/s the cloud peels away below the line: experimental values continue to 0.21 m/s while predictions rarely exceed 0.08 m/s. The model regresses toward the densely sampled low-spread mean.

Panel (b) plots the same predictions as error (predicted minus experimental) against true FSR. At near-zero spread the errors are small and slightly positive: the model overpredicts the slowest flames. As experimental FSR increases, the error trend is a descending line that reaches almost -0.20 m/s at the fastest observations, meaning a high-spread event is predicted near zero. Error variance also widens with true FSR (heteroscedasticity). For a




Figure 4.6: Accuracy–error trade-off for FSR candidates. Mean R^2 increases to the right and mean RMSE decreases downward, so better models lie toward the bottom-right of each panel. Colour denotes model family. Grouped extrapolation (left) and random interpolation (right) show the same family separation—random forest and XGBoost, then KNN, then SVR and MLP—shifted by the protocol. The highlighted points are the selected random-forest baseline.

fire-safety regressor this pattern is the opposite of a conservative bias: the dangerous tail is underpredicted. The hardest held-out paper reaches RMSE 0.1331 m/s, whereas many papers remain below 0.003 m/s. A single pooled R^2 therefore conceals campaign-level failure modes that the error plot makes visible.

4.2.3 FSR explainability


The FSR random forest relies most strongly on thermal diffusion time, oxygen, gas thermal conductivity, and flow velocity in its SHAP summary (figure 4.8). The beeswarm adds direction: high oxygen generally pushes FSR upward, while large thermal diffusion times—a thick or thermally inertial solid—push FSR downward. Held-out permutation importance independently identifies oxygen, flow velocity, gas conductivity, and thermal diffusion time as the largest positive RMSE degradations.

Several correlated material variables have negative permutation importance in the realized held-out sample. That means shuffling happened to reduce error; it does not imply a physically protective effect. The repeated appearance of oxygen, flow, and thermal transport across two explanation methods supports their use as hypotheses and coverage priorities, not as causal coefficients.

4.3 Synthesis


The figures establish one common result across two targets and different model families: within-paper interpolation is strong, while transfer to an unseen paper is moderate and substantially more uncertain. Ignition ROC–AUC falls from approximately 0.91 to 0.72, and PR–AUC from 0.97 to 0.88; FSR RMSE approximately doubles and R^2 falls from 0.879 to 0.458. Family identity is stable: tree ensembles lead both tasks on both protocols. These protocol effects are larger than the differences among the leading tree recipes.

PR–AUC, not ROC–AUC, is the ignition summary that matches a fire-safe screening call, because the label is majority-positive and the operational question concerns the ignition class. For FSR, RMSE and R^2 are two legitimate ranking axes that largely agree, while the residual plot shows that pooled scores hide underprediction of the high-spread tail. The defensible result is a screening tool with explicit domain and uncertainty limits, not a universal flammability law.



(a) Predicted versus experimental FSR on held-out papers. Points below the dashed one-to-one line are underpredictions. (b) Error (predicted minus experimental) versus true FSR. The descending trend is systematic underprediction of fast spread.

Figure 4.7: Grouped random-forest diagnostics. Panel (a) is the extrapolation R^2 view: low-FSR points hug the identity, while high-FSR points fall below it. Panel (b) is the corresponding residual view. Together they show regression toward the low-spread majority and growing error in the sparse high-FSR tail.



(a) Random-forest SHAP values on unseen papers.

(b) Held-out permutation importance.

Figure 4.8: FSR explainability on unseen papers. Thermal diffusion time and oxygen dominate SHAP; permutation importance recovers oxygen, flow, gas conductivity, and thermal diffusion time as the features whose shuffling most increases RMSE.

5 Discussion

The central result of this study is not that one algorithm “solves” material fire safety, but that the apparent predictive difficulty depends strongly on what is held out. When rows were divided at random, experiments from the same publication could occur in both training and test sets. Under that interpolation question, the flame-spread-rate (FSR) random-forest champion reached $R^2 = 0.879 \pm 0.018$ with an RMSE of $0.00626 \pm 0.00047 \text{ m s}^{-1}$, while the ignition XGBoost champion reached a ROC-AUC of 0.908 ± 0.010 . When complete papers were held out, the corresponding results fell to $R^2 = 0.458 \pm 0.222$ and RMSE $0.01221 \pm 0.00412 \text{ m s}^{-1}$ for FSR, and ROC-AUC 0.718 ± 0.044 for ignition. The FSR error approximately doubled, and the ignition ROC-AUC decreased by about 0.19. This consistent row-versus-paper generalization collapse is the most consequential scientific finding. It indicates that models can interpolate effectively among conditions resembling a campaign already represented in training, but only partially transfer to an unseen publication, apparatus, laboratory practice, or sampled regime.

This chapter interprets that finding under four deliberately visible allocations: limitations (approximately 30%), comparison with the literature (approximately 30%), real-world applications (approximately 30%), and future work (approximately 10%). These proportions reflect the practical importance of defining what the models do not establish, placing their scores beside protocol-compatible evidence, and translating them into a cautious decision-support role.

5.1 Limitations

5.1.1 Database heterogeneity and the unit of generalization

The source database is a synthesis of published experiments rather than a prospectively designed factorial study. The FSR analysis contains 2,531 usable target rows from 66 canonical papers; the ignition analysis contains 4,487 deduplicated rows from 87 papers, of which 3,383 report ignition and 1,104 report no ignition. These counts provide considerably more observations than a single campaign, but rows are not independent replications of a common data-generating process. A publication may contribute many nearby settings from one rig, material batch, diagnostic method, gravity platform, or laboratory. Another may contribute only a handful of observations in a distinct regime. Consequently, a row-level split can place near-neighbours—and campaign-specific signatures—on both sides of the evaluation boundary.

The grouped protocol addresses the clearest version of this dependence by constructing canonical paper identities from DOI or citation and prohibiting paper overlap. It therefore tests a harder and more relevant question: transfer to a publication that was absent from model fitting and inner hyperparameter selection. It does not, however, guarantee transfer to every future spacecraft material or operating envelope. Papers can share laboratories, investigators, facilities, wire stocks, reporting conventions, or inherited correlations. Conversely, two campaigns in one paper may differ materially. “Paper” is thus a defensible, reproducible proxy for experimental source, not a perfect ontological definition of domain. A still stronger validation would hold out laboratories, facilities, material families, gravity platforms, or time periods, but the current metadata do not support all such partitions consistently.

The collapse under paper holdout is evidence that publication identity carries predictive structure beyond the intended physical variables. That structure may include legitimate regime information omitted from the schema, such as apparatus heat losses or material morphology, as well as source-specific measurement and reporting practices. The result should not be described simply as overfitting in the narrow algorithmic sense. It is more accurately a combination of source shift, incomplete covariate description, clustered sampling, and possible outcome-definition differences. The large FSR fold variability— R^2 ranging from approximately 0.018 to 0.746 for the grouped champion—further shows that transfer difficulty depends on which papers are withheld.

5.1.2 Missingness, measurement comparability, and representation

Literature synthesis inevitably produces sparse and uneven feature coverage. Material thermophysical properties may be tabulated, inferred from material labels, or absent. Geometry can be reported in formats that are not directly commensurate across wires, flat samples, cylinders, and spheres. Pressure, flow, oxygen concentration, ignition settings, and gravity are not varied independently, and some are fixed or omitted within individual campaigns. The preprocessing contract handles this computationally: numeric values are median-imputed for most families, categorical missingness is assigned an “Unknown” level, and XGBoost can use native numeric missingness. Computational handling is not equivalent to recovering the unobserved physics.

Missingness may itself be source-dependent. If one laboratory systematically reports a property and another does not, a tree can use the missingness pattern as a source fingerprint. This can improve random-split prediction while harming transfer. Median imputation also compresses genuinely different unknown cases to a common value, while a constant categorical level conflates “not measured”, “not applicable”, and “not reported”. The

grouped results therefore integrate both the difficulty of physical extrapolation and the consequences of imperfect harmonisation.

Feature engineering improves consistency but does not remove these concerns. The ignition pipeline uses 32 numeric and five categorical features, including material, gas-transport, operating-condition, geometry, flow-direction, gravity-regime, and ignition-method descriptors. The FSR pipeline detects the available numeric and categorical roles after excluding explicit source identifiers, duplicate target fields, and post-outcome variables. Derived quantities such as Reynolds, Peclet, Prandtl, volumetric heat capacity, thermal diffusivity, and thermal diffusion time can encode useful physical structure. Yet derived features share inputs and can be mutually dependent; their values are only as reliable as the source measurements and assumptions used to calculate them.

One schema discrepancy requires explicit qualification: the persisted FSR feature manifest includes categorical **Ignition**, although the pipeline documentation describes ignition fields as excluded. Its reported tree and SHAP importance is zero, but that does not substitute for a controlled ablation. Future FSR evaluation should remove the field and regenerate frozen comparisons before asserting that the documented and executed schemas are identical.

5.1.3 Imbalance and observational sampling

Two forms of imbalance matter. First, ignition is a majority-positive task: approximately three quarters of the 4,487 records are ignition cases. A high PR-AUC is therefore partly facilitated by prevalence and must be read beside ROC-AUC, Brier score, sensitivity, specificity, and threshold-dependent performance. Second, papers contribute unequal row counts. A densely sampled campaign can dominate empirical risk even if it represents a narrow region. The candidate portfolios test no, square-root, and inverse paper weights, and weight-blind families use deterministic probability-proportional resampling. These mechanisms change each source’s influence, but cannot manufacture missing regimes. The selected grouped champions were the unweighted baselines: random forest for FSR and XGBoost for ignition. This does not prove that source imbalance is unimportant. It only shows that the tested weighting corrections did not improve the selection objective enough under these frozen splits.

Observational coverage creates a more fundamental limitation. Oxygen, pressure, gravity, flow, geometry, material properties, and ignition practice were chosen by investigators for particular scientific questions. Their joint distribution is therefore neither random nor uniform. For example, a material may occur mainly in one geometry and gravity regime, while a pressure range may be associated with one facility. A flexible learner can predict from these correlations without separating their physical effects. Sparse partial-gravity observations are especially problematic because lunar and Martian gravity conditions cannot be treated as densely sampled interpolation domains merely because gravity is a numeric feature.

5.1.4 Threshold policies and the operating points they induce

A ROC or PR curve ranks cases; a deployed screener still needs a number τ at which a probability becomes a yes. Four policies were frozen from inner out-of-fold scores: maximum MCC, maximum F1, maximum balanced accuracy, and maximum Youden J . Balanced accuracy and Youden J selected the same thresholds in both champions, so figure 5.1 shows three distinct policies for each of the interpolation and extrapolation models.

The figure plots sensitivity (recall) against specificity. The dashed diagonal is the locus of equal sensitivity and specificity; points above it catch ignitions at a higher rate than they correctly clear non-ignitions, and points below it do the reverse. Marker shape distinguishes the interpolation champion (**xgb_paper_sqrt**, circles) from the extrapolation champion (**xgb_baseline**, triangles). Colour distinguishes the policy.

Interpolation occupies the upper-right quadrant. At $\tau \approx 0.84$ the balanced-accuracy policy attains sensitivity 0.77 and specificity 0.90; MCC at $\tau \approx 0.79$ is similar (0.79, 0.87). Those high thresholds are usable screening settings inside known campaigns: most ignitions are still caught, and most true non-ignitions remain unlabeled. The F1 policy drops the threshold to $\tau \approx 0.37$ and induces a different operating point: sensitivity 0.95 but specificity only 0.48. Because F1 ignores true negatives, it buys almost-complete recall by labeling many non-ignitions as ignition. Even within interpolation, F1 is therefore not a fire-safe default; it is a high-recall policy whose false-alarm cost must be accepted explicitly.

Extrapolation induces a much harsher map. All three triangles lie on or below the diagonal. Balanced accuracy at $\tau \approx 0.86$ sits on the diagonal at (0.66, 0.66): the grouped champion, once forced to a single cut, is no better than a balanced coin in this sensitivity-specificity plane. MCC at $\tau \approx 0.81$ trades a little more recall (0.73) for weaker specificity (0.58). F1 at $\tau \approx 0.29$ produces sensitivity 0.98 and specificity 0.06: the model essentially predicts ignition for every held-out case. A “never miss a fire” threshold on a new paper therefore does not yield a selective screener; it yields a near-constant positive alarm.

Two consequences follow. First, no threshold trained on interpolation transfers as a universal operating point: the same policy names produce qualitatively different (sensitivity, specificity) pairs once papers are withheld. Second, metric choice is itself a safety decision. Optimizing F1, or any recall-heavy scalar, induces an alarm-almost-always rule under source shift, which would saturate a test campaign rather than prioritize it. Optimizing balanced accuracy or MCC induces a more even trade-off, but under extrapolation that trade-off is only moderate and is not a certification margin. The curves in figure 4.1 remain valid rankings; figure 5.1 shows that converting those rankings into a binary fire-safe call still requires an explicit loss function, and that the grouped evidence does not currently support a high-recall, high-specificity cut.




Figure 5.1: Ignition operating points induced by three threshold policies for the interpolation champion (circles) and the extrapolation champion (triangles). Axes are sensitivity (recall) and specificity. F1 selects a low threshold and collapses specificity, especially under paper hold-out ($\tau \approx 0.29$, specificity 0.06). Balanced accuracy and MCC select high thresholds that remain usable in interpolation but only moderately balanced—or worse—once entire papers are withheld.

5.1.5 Model dependence, metrics, and uncertainty

Model rankings are contingent on the tested candidates, search spaces, preprocessing, loss functions, seeds, and selection policy. The configuration files declare 34 FSR and 26 ignition candidates, but the persisted merged FSR manifest contains 25 candidates and the formal ignition selection record omits the available MLP shards. The selected FSR champion was random forest rather than XGBoost, while ignition selected XGBoost. These results apply to the persisted comparison pools; they neither establish complete execution of every configured recipe nor prove that the selected model is optimal among all possible algorithms.

Likewise, a point estimate is not a complete account of performance. R^2 is sensitive to the target variance in a held-out fold and can be unstable when a small set of papers defines a narrow or shifted range. FSR MAPE is extremely large because percentage error is ill-conditioned when observed spread approaches zero; RMSE and MAE in physical units are more interpretable here. For ignition, ROC-AUC measures ranking across thresholds but not probability calibration or a mission-specific operating point. The paper-cluster bootstrap, repeated splits, per-paper metrics, calibration outputs, and fold standard deviations should therefore accompany any headline number.

The final refitted artifacts also have a different evidential status from the outer-fold results. Refitting uses all labelled rows and is appropriate for producing a deployable model, but its training fit is not an unbiased estimate of future performance. Only the nested, untouched outer-fold predictions support comparative claims. The model cards correctly preserve this distinction.

5.1.6 SHAP and the boundary between explanation and causation

The explainability outputs describe model behaviour, not causal mechanisms. For ignition, the global XGBoost SHAP ranking assigns high mean absolute contributions to oxygen, absolute flow speed, Prandtl number, gas molar mass, fuel pyrolysis temperature, and ignition time. Yet grouped permutation importance places fuel pyrolysis temperature, ignition time, fuel heat capacity, thermal diffusion time, and a sample dimension at the top. For FSR, SHAP ranks thermal diffusion time and oxygen prominently, while held-out permutation importance ranks oxygen, flow velocity, gas thermal conductivity, and thermal diffusion time most strongly. Differences are expected because SHAP and permutation importance answer different questions and react differently to correlated inputs.

A SHAP value is an additive decomposition of a particular fitted function relative to a background distribution. It does not show that intervening on that feature would produce the attributed change. Correlated and derived variables allow credit to be redistributed among substitutes. Source confounding can make an apparatus-correlated variable appear physically decisive. Missingness patterns may attract attribution. Mean absolute SHAP also suppresses direction, interactions, and regime dependence. The physically plausible appearance of many rankings is reassuring, but plausibility is not validation of a causal law. Claims such as “oxygen causes the predicted increase” require intervention, mechanistic identification, or a causal design that this observational database does not provide.

The safest interpretation is predictive: within the fitted model and sampled distribution, these features are influential for separating or estimating outcomes. The permutation results add a transfer-oriented check by measuring score degradation on held-out data, but they remain model- and sample-dependent. Explainability is therefore best used to audit whether the learner relies on obviously inappropriate signals, to generate hypotheses, and to identify data-collection priorities. It should not replace combustion theory or controlled experimentation.

5.2 Comparison with the Literature

5.2.1 Rivera/Jose wire-specific machine learning

Rivera, Jose, and colleagues report a literature-derived electrical-wire FSR study in which categorical regime descriptors and augmentation improve random ten-fold cross-validation performance, with several nonlinear models exceeding $R^2 = 0.9$ [9]. Their database contains approximately 1,200 wire observations and explicitly represents core and insulation materials, spread direction, and configuration. Categorisation raises the reported performance of gradient boosting, KNN, SVM, decision trees, and MLP, while bootstrap augmentation is particularly effective for gradient boosting and KNN. This supports two conclusions also visible here: flame spread is strongly nonlinear, and categorical information can expose regime structure that numeric properties alone may not capture.

The numerical results must not be presented as if they were generated by equivalent protocols. The Rivera/Jose work uses random cross-validation and explicitly characterises its result as robustness within the compiled database rather than strict leave-source-out transfer. Its scope is electrical wires, and its target is reported in mm s^{-1} . The present FSR database is broader in geometry and source coverage, reports the model target in m s^{-1} , and makes paper-disjoint nested evaluation the primary scientific test. Hyperparameter search, feature definitions, data versions, augmentation, and source composition also differ. A direct statement that 0.458 is “worse than” > 0.9 would confound task and protocol.

The fair comparison is question-specific. On row-level interpolation, this study’s random-forest $R^2 = 0.879 \pm 0.018$ is close to, though generally below, the strongest wire-specific random-CV values. On paper-level extrapolation, the same family falls to $R^2 = 0.458 \pm 0.222$. Rivera/Jose do not report an equivalent paper-disjoint score,

so their > 0.9 result neither contradicts nor establishes performance on this harder transfer task. Instead, their stated caveat and the present protocol gap jointly reinforce the need to name the validation unit whenever quoting accuracy.

There is also a substantive difference in what categorisation can encode. Wire-specific labels such as insulation chemistry, core material, concurrent/opposed spread, and orientation can partition physically distinct regimes. In a broader database, adding similarly detailed categories may improve prediction, but it may also increase sparsity and source fingerprinting. The present models exclude explicit paper identifiers and post-outcome features, yet legitimate categories can remain highly campaign-specific. Their value should therefore be tested under grouped validation, not assumed from random-CV gains.

The augmentation comparison is similarly instructive. Rivera/Jose find bootstrap resampling particularly useful, while noise injection and GAN augmentation can create statistically plausible but not necessarily physically consistent combinations. The present FSR benchmark treats bootstrap copies as reweighting observed conditions; copies inherit the source paper and are never added to the test set. Its selected FSR champions use the real-data condition rather than bootstrap augmentation. Thus, augmentation did not compensate for source shift in the selected model. This supports a conservative principle: resampling can alter optimisation and density, but it cannot supply a new material, apparatus, or gravity regime.

5.2.2 XGBoost and the tabular-learning literature

The strength of tree ensembles is consistent with broader tabular-learning evidence. Schwartz-Ziv and Armon compare XGBoost with several deep tabular architectures on eleven datasets and find that XGBoost generally outperforms individual deep models, requires less tuning, and is less sensitive to whether a dataset appeared in the proposing model’s original paper [11]. The present database has precisely the characteristics that often favour tree ensembles: mixed numeric and categorical predictors, nonlinear interactions, sparse cells, missing fields, unequal scales, and modest sample size relative to the variety of regimes.

The ignition outcome fits this expectation strongly. XGBoost occupies the top of both interpolation and grouped leaderboards; the paper-square-root variant is selected for interpolation and the unweighted baseline for extrapolation. By comparison, decision tree, KNN, and SVM candidates generally rank lower, while the tested focal XGBoost variants perform near chance. This last result is a reminder that a generally strong architecture does not rescue an unsuitable objective implementation or weighting combination.

The FSR result is more nuanced. Random forest narrowly leads interpolation and is selected for grouped transfer, while several XGBoost variants remain competitive. The selected random forest uses no paper weighting, no augmentation, and standard identity-target regression. Its grouped R^2 is 0.458, compared with 0.441 for the strongest listed augmented paper-weighted XGBoost candidate. This difference is small relative to fold variation and should not be elevated into a universal claim that bagging is superior to boosting. It shows only that, within the tested portfolio and policy, the random-forest baseline offered the preferred accuracy–stability combination.

The MLP candidates do not outperform the tree ensembles, particularly under grouped FSR evaluation, where their mean R^2 values can become negative and highly variable. This agrees directionally with the tabular literature, but the explanation is not that neural networks are intrinsically inappropriate. Data volume, regularisation, architecture, optimisation, and representation all matter. Schwartz-Ziv and Armon also find that ensembles combining deep models with XGBoost can outperform either component alone [11]. The present work did not evaluate a formally stacked, outer-fold-safe heterogeneous ensemble, so it cannot rule out such gains.

5.2.3 Physics-informed learning and combustion knowledge

Physics-informed neural networks constrain learning with residuals of governing differential equations and boundary or initial conditions [8]. More broadly, physics-informed machine learning seeks to combine data with conservation laws, symmetries, constitutive knowledge, or admissibility constraints [4]. This is attractive for flame spread because experimental coverage is sparse, predictions must remain physically plausible, and extrapolation is precisely where unconstrained correlations are weakest.

The present models are physics-featured rather than physics-constrained. Inputs include thermophysical properties, dimensionless groups, geometry, oxygen, pressure, flow, gravity, and ignition conditions, but the learned function is not required to satisfy mass, species, momentum, or energy balances. Optional monotonic constraints on oxygen exist for selected XGBoost configurations, yet a monotonic feature constraint is not equivalent to a combustion model. Likewise, categorising gravity or spread direction supplies domain structure but does not create a PINN.

Applying a conventional PINN directly would require an agreed governing formulation, state fields, boundary conditions, and sufficient spatial or temporal observations. This database predominantly contains tabular experimental summaries and endpoint outcomes, not common-domain field solutions. The Rivera/Jose paper also notes that the theoretical framework remains incomplete for direct physics-informed deployment [9]. Therefore, the immediate opportunity is more modest: enforce nonnegative FSR, use dimensionally consistent representations, penalise known inadmissible combinations, model regime-specific monotonicity where defensible, or hybridise reduced-order heat-transfer relations with learned residuals. Such constraints must still be evaluated on held-out sources; physical plausibility alone does not guarantee empirical transfer.

5.2.4 Spacecraft fire-safety context

Guibaud and colleagues identify fire as a major concern for inhabited exploration, particularly as missions move farther from immediate Earth support and involve different atmospheres, gravity levels, confinement, and operational constraints [2]. In that setting, ignition and flame spread are not abstract prediction targets. They contribute to material selection, wire and component design, experiment prioritisation, contingency planning, and the evaluation of how a local ignition might develop.

The present findings add a data-governance qualification to that motivation. Historical evidence can support useful interpolation, but mission-relevant environments may resemble a new source domain rather than another row from a known campaign. Lunar and Martian partial gravity, altered oxygen and pressure, aging materials, realistic ventilation, and integrated spacecraft geometries may lie in sparsely observed combinations. The correct literature-level conclusion is therefore not that machine learning replaces fire testing, but that it can organise heterogeneous evidence and identify where that evidence is or is not transferable.

5.3 Real-World Applications (approximately 30%)

5.3.1 Screening and prioritisation

The most defensible near-term application is screening. Given candidate conditions in the documented schema, the ignition artifact can rank cases by predicted ignition probability, and the FSR artifact can estimate spread rate for records sufficiently similar to the training support. This can reduce a large design space to a smaller set requiring detailed simulation or experiment. A screening tool can, for example, flag material–geometry–atmosphere combinations whose predictions and uncertainty suggest elevated concern, or prioritise boundary cases where models disagree.

Such use should be explicitly asymmetric. A high-risk prediction may justify additional investigation; a low-risk prediction should not by itself waive testing. False reassurance is particularly consequential in spacecraft fire safety. Threshold selection must therefore follow an operational loss function rather than a generic metric. If missing a potentially ignitable case is costly, a sensitivity-oriented threshold may be appropriate, accepting more false alarms. If scarce test capacity is the binding constraint, a more specific threshold may be chosen, but the resulting missed-case risk must be quantified. Four policies—MCC, F1, balanced accuracy, and Youden J —are retained rather than declaring one universally correct; figure 5.1 shows that F1 in particular induces an almost-always-positive rule under paper hold-out.

FSR screening should rely primarily on RMSE and MAE in physical units, not MAPE near zero. Predictions should be accompanied by a domain flag and an interval or empirical error summary. The grouped champion’s mean RMSE of 0.0122 ms^{-1} , with substantial fold variation, is large relative to many low-spread observations. Ranking may therefore be more reliable than precise point estimation in underrepresented regimes. Where a decision depends on a narrow FSR boundary, confirmatory testing is necessary.

5.3.2 Experiment design and the unexplored oxygen–pressure domain

A second application is directing new data collection. Per-paper errors can identify campaigns that are systematically difficult to predict. Disagreement between random and grouped performance identifies fields that may encode source-specific rather than transferable information. Explainability, used cautiously, can nominate influential variables whose measurement completeness deserves improvement. For example, the repeated importance of fuel pyrolysis temperature, ignition time, thermal diffusion time, oxygen, flow, and gas properties suggests that missing or inconsistently defined values in these fields can be costly.

The most direct map of what remains to be tested is the oxygen–pressure coverage of the assembled table (figure 5.2). Panel (a) plots every observation in oxygen mole fraction versus pressure, coloured by ignition outcome. Panel (b) bins the same plane by sampling density; empty cells are empirical gaps, not physically forbidden regions. Most rows cluster in a horizontal band near terrestrial pressure ($\approx 100 \text{ kPa}$) and in vertical stripes at a few laboratory oxygen set-points (0.2, 0.3, 0.4, 0.5, and 1.0). Ignition and no-ignition outcomes overlap inside that band: neither oxygen nor pressure alone separates the label, which is why a univariate oxygen cutoff is not a substitute for the multivariate model. Equally important are the white regions. Oxygen mole fractions between about 0.55 and 0.95 are almost unsampled except at isolated pressures; low oxygen below 0.1 is nearly empty; high pressure above a few megapascals is visited only at a handful of oxygen fractions; and low pressure below $\approx 20 \text{ kPa}$ is sparse away from a small low-oxygen tail. Proposed exploration atmospheres—reduced total pressure with elevated oxygen fraction—sit precisely in combinations that the density plot leaves blank or thinly populated.

This geometry is the practical content of “out of distribution” for a spacecraft material screener. A query whose oxygen and pressure each lie inside the marginal ranges can still fall in an empty joint cell. Predictions there are extrapolations in the experimental design sense, even if the model API returns a number. The figure therefore does not merely decorate the database: it is a search-space map. Future tests should be allocated to those empty oxygen–pressure cells, crossed with underrepresented geometries, materials, and gravity platforms, rather

than added as further dense repeats of the 100 kPa band. Because spaceflight experiments are expensive, model uncertainty and this coverage plot can help allocate them, but the acquisition strategy should reserve samples for prospective validation rather than using every new result for immediate retraining.

A proposed test condition can be compared with ranges, categories, and combinations represented in training. A case in a familiar oxygen–pressure range but a new material–geometry–gravity combination should still be labelled out of distribution. Partial-gravity testing remains an especially strong candidate for prioritisation: lunar- or Martian-relevant gravity, with controlled atmosphere and flow, would test whether a numeric gravity feature transfers smoothly or whether regime-specific dynamics dominate.




Figure 5.2: Oxygen mole fraction versus pressure in the assembled database (logarithmic pressure axis). The dense terrestrial-pressure band and a few discrete oxygen set-points dominate the literature. Empty regions—high oxygen at intermediate pressure, very low oxygen, and much of the high-pressure and low-pressure envelopes—are the experimental search space that screening still cannot claim.

5.3.3 Material and design decision support

Within a formal engineering workflow, predictions can be one line of evidence in a multi-criteria decision analysis. Candidate materials may be evaluated on ignition propensity and expected spread alongside toxicity, smoke, mechanical performance, electrical function, mass, manufacturability, aging, and compatibility. The models can reveal trade-offs: one material may score lower for ignition under represented conditions yet have uncertain spread behaviour in a new geometry. Presenting both tasks separately is essential because ignition probability and post-ignition spread rate are related but not interchangeable safety endpoints.

A suitable report for each candidate would contain: the input record and units; missing or imputed fields; the model and data fingerprint; interpolation and grouped evidence; the selected threshold policy; predicted probability or FSR; uncertainty or fold-based error context; nearest supported regimes; out-of-distribution warnings; and known limitations. It should also record whether the recommendation is merely “test first”, “requires expert review”, or “within represented screening envelope”. Raw probabilities without this provenance invite overinterpretation.

Human review remains central. A combustion specialist should assess whether the implied trend is physically credible and whether the case violates assumptions hidden by the tabular representation. A materials or systems engineer should evaluate integration effects absent from the database. A safety authority should determine the acceptable evidence and margins. Model output is an aid to these roles, not an automated disposition.

5.3.4 Uncertainty, abstention, and out-of-distribution use

The current repeated-fold standard deviations and cluster-bootstrap intervals provide evaluation uncertainty, not a complete case-specific predictive distribution. Operational use needs both aleatoric uncertainty, arising from experimental variability, and epistemic uncertainty, arising from sparse support and model uncertainty. An ensemble across seeds, folds, or model families could estimate disagreement. Conformal methods could provide empirical intervals or prediction sets if calibrated under a defensible exchangeability assumption. Neither removes distribution shift; grouped or source-conditional calibration remains necessary.

Abstention is a safety feature. The system should decline to issue a routine prediction when required fields are absent, categorical levels are unseen, the query is distant from training support, model portfolios disagree substantially, or a physically admissible range is violated. Out-of-distribution detection should consider mixed

data types and joint structure. A simple min–max range check is insufficient because impossible or unsupported combinations can lie inside all marginal ranges.

For ignition, calibration must be examined under held-out papers. A nominal probability of 0.8 is operationally meaningful only if similar grouped predictions ignite at approximately that rate, within uncertainty. For FSR, an interval should be checked for coverage by paper and regime, not only pooled rows. Performance dashboards should expose the worst papers and regimes rather than average away failure pockets.

5.3.5 No replacement for certification or qualification

These artifacts are not certification models. They have not undergone prospective external validation, formal software qualification, independent verification and validation, uncertainty qualification for a specified design envelope, or acceptance by a regulatory or mission assurance authority. The database is observational and heterogeneous; source-held-out performance is moderate; and the models do not enforce governing physics. No model output should replace required material flammability tests, wire qualification, hazard analysis, engineering judgement, or certification evidence.

This limitation does not make the models unusable. Screening and decision support can be valuable precisely because they precede authoritative testing. The proper claim is that the models summarise existing evidence and can prioritise resources under documented uncertainty. The improper claim is that a predicted “no ignition” or low FSR establishes compliance.

Deployment would require version control and auditability. Data hashes, feature manifests, split assignments, selected hyperparameters, model cards, and out-of-fold predictions are already recorded. An operational service should additionally lock the software environment, validate units and schemas, log every query and response, monitor drift, and trigger review when the database or artifact changes. Evaluation artifacts and refitted artifacts must remain clearly separated: the former support performance claims; the latter generate predictions.

5.3.6 A staged deployment concept

A cautious implementation can proceed in stages. In an offline research stage, users inspect leaderboards, paper-level errors, learning curves, calibration, and explanations without affecting decisions. In a shadow stage, predictions accompany conventional assessments but are not used to change them; discrepancies are adjudicated and logged. In a controlled screening stage, outputs may prioritise tests within a predeclared envelope, with mandatory abstention rules and expert sign-off. Only after external prospective evidence could a broader role be considered.

Success criteria should be decision-centred. A useful screener is not merely one with a high ROC–AUC; it reduces testing burden or identifies hazardous cases without exceeding a tolerated miss rate, while remaining calibrated across sources. A useful regressor does not merely minimise pooled RMSE; it provides intervals that cover new-paper observations and rankings that support experimental prioritisation. These criteria align model development with fire-safety outcomes rather than leaderboard optimisation.

5.4 Future Work (approximately 10%)

The immediate priority is external, prospective validation. New experiments should be registered before outcomes are observed, processed through a frozen artifact, and retained as an untouched external test set. Validation should report both pooled and source-specific performance, calibration, abstention frequency, and errors by material, geometry, atmosphere, and gravity regime. Only after this evaluation should the new records enter training.

More partial-gravity data are needed. Purposeful lunar- and Martian-gravity campaigns should vary oxygen, pressure, flow, material, and geometry sufficiently to distinguish gravity effects from source and apparatus effects. Replicate measurements and harmonised uncertainty reporting would separate experimental variability from domain shift.

Uncertainty and out-of-distribution detection should become first-class outputs. Paper-aware ensembles, grouped conformal calibration where assumptions are supportable, mixed-type distance or density models, and explicit abstention thresholds should be evaluated prospectively. Coverage should be measured by source and regime, not only globally.

The database would benefit from a shared ontology: canonical material identifiers; geometry and orientation definitions; units; uncertainty fields; missingness reasons; facility, rig, and diagnostic metadata; and controlled vocabularies for ignition and extinction outcomes. Better provenance would allow laboratory- or facility-held-out validation and reduce accidental source fingerprints.

Finally, physically constrained models should be tested without assuming that a full PINN is immediately appropriate. Useful steps include nonnegative FSR outputs, dimensionally consistent features, regime-aware monotonic constraints, conservation-inspired penalties, and hybrids that learn corrections to reduced-order flame-spread relations. Fully physics-informed formulations become credible only when governing equations, boundary conditions, and compatible field data are available [4, 8]. Every proposed constraint must be assessed under the same paper-disjoint and prospective protocols used to test statistical models.

6 Conclusion

This research asked whether heterogeneous literature data can support reliable machine-learning prediction of material ignition and flame spread rate for fire-safe spacecraft design, especially when prediction is required for an experimental source that the model has not previously seen. The precise answer is conditional. The database supports strong interpolation among rows from represented papers, but only moderate and substantially less stable generalisation to held-out papers. It is therefore suitable as the basis of a traceable screening and research-prioritisation tool, but it does not yet support universal prediction, causal claims, or replacement of qualification and certification testing.

For FSR regression, 2,531 usable observations from 66 canonical papers were evaluated with frozen random-row and paper-grouped holdouts. The selected random-forest baseline achieved $R^2 = 0.879 \pm 0.018$, RMSE = $0.00626 \pm 0.00047 \text{ m s}^{-1}$, and MAE = $0.00210 \pm 0.00017 \text{ m s}^{-1}$ under random-row interpolation. When complete papers were withheld, the same model achieved $R^2 = 0.458 \pm 0.222$, RMSE = $0.01221 \pm 0.00412 \text{ m s}^{-1}$, and MAE = $0.00482 \pm 0.00114 \text{ m s}^{-1}$. For ignition classification, 4,487 deduplicated observations from 87 papers were evaluated. The paper-square-root-weighted XGBoost interpolation champion achieved ROC-AUC = 0.908 ± 0.010 and PR-AUC = 0.969 ± 0.004 ; the unweighted XGBoost grouped champion achieved ROC-AUC = 0.718 ± 0.044 , PR-AUC = 0.877 ± 0.030 , and Brier score = 0.179 ± 0.017 .

The gap between these protocols is the primary result. Random cross-validation answers whether a model can recover nearby conditions in campaigns it has partly observed. Paper-grouped validation asks whether it can transfer across a publication boundary that bundles apparatus, laboratory practice, sampled regime, and reporting convention. The latter is the more relevant internal evidence for a new campaign, although it is not a substitute for external validation. Quoting only the random-split scores would materially overstate readiness.

The model comparison also answers the algorithmic part of the question. Tree ensembles are well matched to these mixed, sparse, nonlinear tabular data. XGBoost led ignition prediction, while random forest was selected for FSR; KNN, SVM/SVR, decision-tree, and MLP candidates supplied useful baselines but did not displace the grouped champions. This is consistent with tabular-learning literature, while also demonstrating that XGBoost is not universally best [11]. The wire-specific $R^2 > 0.9$ findings of Rivera/Jose and colleagues remain compatible with the present results because they concern random cross-validation within a narrower electrical-wire database, not paper-disjoint transfer [9].

Feature attributions identified physically plausible variables, including oxygen, flow, thermal diffusion time, fuel pyrolysis temperature, ignition time, and gas or material properties. These findings are diagnostic rather than causal. SHAP decomposes the fitted model; it does not identify the effect of an intervention. Correlated derived variables, uneven missingness, and campaign confounding prevent causal interpretation.

Accordingly, the practical contribution is a reproducible decision-support framework. It freezes splits, isolates tuning within outer training data, verifies paper disjointness, records data fingerprints and feature manifests, reports repeated-fold and paper-level uncertainty, and distinguishes unbiased evaluation from full-data refitting. Within a documented applicability domain, its predictions can rank candidate tests, identify poorly covered regimes, and focus expert attention. They must be accompanied by uncertainty, out-of-distribution checks, threshold rationale, and explicit abstention.

The work therefore establishes a useful but bounded capability: literature-scale machine learning can extract transferable signal for spacecraft fire-safety screening, yet database heterogeneity and source shift remain the dominant obstacles. The path to stronger evidence is not a larger headline score on mixed rows. It is prospective external validation, targeted partial-gravity experiments, a harmonised ontology, uncertainty-aware abstention, and physically constrained models evaluated against genuinely new sources. Until those steps are completed, the system should inform experiments and engineering judgement, never replace them.

A Technical Appendices

A.1 Feature dictionary

The deployed manifests are the authoritative machine-readable feature definitions. Tables A.1 and A.2 provide a compact human-readable dictionary. Units are encoded in names where available; otherwise they follow the database schema. Derived quantities are predictors, not independent measurements.

Table A.1: Ignition feature dictionary (**fable-data-v5**; 32 numeric and five categorical features).

Group	Features and interpretation
Fuel properties	<code>fuel_density_kg_m3</code> , <code>fuel_k_w_mk</code> , <code>fuel_cp_j_kgk</code> , <code>fuel_pyrolysis_t_k</code> , <code>fuel_alpha_m2_s</code> , and <code>fuel_volumetric_heat_capacity_j_m3k</code> : density, conductivity, specific heat, pyrolysis temperature, diffusivity, and volumetric heat capacity.
Core properties	<code>core_density_kg_m3</code> , <code>core_k_w_mk</code> , <code>core_cp_j_kgk</code> , and <code>core_volumetric_heat_capacity_j_m3k</code> .
Gas and transport	<code>oxygen</code> , <code>gas_molar_mass</code> , <code>gas_cp_j_kgk</code> , <code>gas_k_w_mk</code> , <code>gas_density_kg_m3</code> , <code>gas_alpha_m2_s</code> , <code>gas_nu_m2_s</code> , <code>reynolds</code> , <code>peclet</code> , <code>prandtl</code> , and <code>thermal_diffusion_time_s</code> .
Operating conditions	<code>pressure_pa</code> , <code>signed_flow_velocity_m_s</code> , <code>flow_speed_abs_m_s</code> , <code>gravity_g</code> , <code>ignition_power_w</code> , and <code>ignition_time_s</code> .
Geometry	<code>half_thickness_m</code> , <code>characteristic_length_m</code> , and <code>sample_dim_1_m</code> through <code>sample_dim_3_m</code> .
Categorical regime	<code>geometry_cat</code> (wire, flat, cylindrical, spherical), <code>diluent_cat</code> , <code>flow_direction_cat</code> (coflow, counterflow, quiescent), <code>gravity_regime_cat</code> (microgravity, partial, Earth, hyper/unknown), and <code>ignition_method_cat</code> .

Table A.2: FSR feature dictionary (**fsr-data-v2**; exact availability follows the artifact manifest).

Group	Manifest fields
Geometry	Sample dimensions, <code>half_thickness</code> , <code>characteristic_length_m</code> , and categorical <code>Geometry of Sample</code> .
Fuel properties	<code>fuel_density_kg_m3</code> , <code>fuel_k_W_mK</code> , <code>fuel_cp_J_kgK</code> , <code>fuel_pyrolysis_T_K</code> , <code>fuel_alpha_m2_s</code> , and <code>fuel_volumetric_heat_capacity_J_m3K</code> .
Core properties	<code>core_density_kg_m3</code> , <code>core_k_W_mK</code> , <code>core_cp_J_kgK</code> , and <code>core_volumetric_heat_capacity_J_m3K</code> .
Gas and transport	Oxygen Concentration, <code>gas_M</code> , <code>gas_cp_mass</code> , <code>gas_k</code> , <code>gas_density_kg_m3</code> , <code>gas_alpha_m2_s</code> , <code>gas_nu_m2_s</code> , <code>Reynolds</code> , <code>Peclet</code> , <code>Prandtl</code> , and <code>thermal_diffusion_time</code> .
Conditions and apparatus	Pressure (Pa), Flow Velocity (m/s), Gravity, Ignition Power (W), Ignition Time (s), categorical diluent, Ignition Method, and antecedent categorical <code>Ignition</code> .

Explicit source identifiers, DOI/citation fields, duplicate FSR targets, and post-outcome variables such as flame length, heat-release rate, smoke, and extinction indicators are excluded from predictors. Paper identity is retained only for grouping and weighting. Numeric parsing accepts cells beginning with a number; strings such as N2 and CO2 remain categorical. For non-XGBoost models, numeric missing values are median-imputed and categorical values receive a constant unknown level; XGBoost may retain native numeric missingness.

A.2 Metric definitions

For observed FSR y_i , prediction $\hat{y}_i$, sample mean $\bar{y}$, and n test observations,

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}, \quad (\text{A.1})$$

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}, \quad (\text{A.2})$$

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|, \quad (\text{A.3})$$

$$\text{MBE} = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i), \quad (\text{A.4})$$

$$\text{MAPE} = \frac{100}{n} \sum_{i=1}^n \left| \frac{y_i - \hat{y}_i}{y_i} \right|. \quad (\text{A.5})$$

NRMSE is a normalised RMSE. MAPE is unstable when y_i approaches zero and is secondary to RMSE and MAE for this FSR distribution.

For ignition probability $\hat{p}_i$, the Brier score is

$$\text{Brier} = \frac{1}{n} \sum_{i=1}^n (\hat{p}_i - y_i)^2. \quad (\text{A.6})$$

ROC–AUC measures ranking of positives above negatives across thresholds. PR–AUC summarises precision–recall performance and must be read in light of class prevalence. Given threshold τ ,

$$\hat{y}_i = \mathbf{1}[\hat{p}_i \geq \tau]. \quad (\text{A.7})$$

Sensitivity is $\text{TP}/(\text{TP} + \text{FN})$, specificity is $\text{TN}/(\text{TN} + \text{FP})$, precision is $\text{TP}/(\text{TP} + \text{FP})$, and balanced accuracy is the mean of sensitivity and specificity. The F1 score is the harmonic mean

$$\text{F1} = \frac{2 \text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}. \quad (\text{A.8})$$

The Matthews correlation coefficient is

$$\text{MCC} = \frac{\text{TPTN} - \text{FPFN}}{\sqrt{(\text{TP} + \text{FP})(\text{TP} + \text{FN})(\text{TN} + \text{FP})(\text{TN} + \text{FN})}}. \quad (\text{A.9})$$

Thresholds optimising MCC, F1, balanced accuracy, and Youden’s J are selected from inner out-of-fold predictions, never from the corresponding outer test labels.

A.3 Reproducibility workflow

The reproducibility contract is:

1. Load the named database and record its SHA-256 fingerprint, data version, parsing decisions, target counts, missingness, and duplicate removal.
2. Canonicalise DOI or citation into **paper_id**; exclude source identifiers and post-outcome leakage from model inputs.
3. Freeze all row and paper-grouped split assignments before model comparison. Assert complete row coverage, no row overlap, and zero paper overlap for grouped protocols.
4. For every candidate and outer fold, fit preprocessing only on the outer training partition. Tune hyperparameters on inner folds; use group-disjoint inner folds when the outer protocol is grouped.
5. Fit the selected inner configuration on the complete outer training partition and score the untouched outer test partition once. For ignition, freeze thresholds from inner out-of-fold probabilities. For FSR bootstrap conditions, augment training only and preserve inherited paper identity.
6. Persist outer predictions, fold and per-paper metrics, search histories, selected hyperparameters, integrity checks, and failure records. Aggregate only complete, valid candidates.
7. Select champions from persisted tables under the versioned policy. The FSR policy prioritises RMSE; the ignition policy prioritises ROC–AUC and requires LOPO robustness for extrapolation.

8. Refit the selected configuration on all labelled data for inference, package the model card and feature manifest, and retain evaluation results separately. A refitted artifact is deployable; outer-fold predictions are the unbiased comparison evidence.

The recorded fingerprints are shown below; random seeds begin at 42 and split files are shared across candidates.

FSR: 6e09a99e4ea69148e69151bd46149be4140d181b2b8294106927ea9499002fc8
 Ignition: d80da0648ca3a7d77d31ae90e53b4e81b2ab28c8ec666ce6cd150ee7544f6

A.4 Traceability matrix

Table A.3: Claim-to-evidence traceability matrix.

Claim	Persisted evidence
FSR dataset has 2,531 target rows from 66 papers	Dataset-integrity summary for the FSR modeling cohort.
FSR interpolation and grouped champion scores	Frozen interpolation and extrapolation leaderboards and the FSR selection record.
FSR transfer loss and stability	Repeated-holdout gap and metric-stability tables.
Ignition dataset has 4,487 rows, 87 papers, and 3,383/1,104 class counts	Dataset-integrity summary for the ignition modeling cohort.
Ignition interpolation and grouped scores	Frozen interpolation and extrapolation leaderboards.
Threshold trade-offs	Champion threshold-performance table used in figure 5.1.
Explanations are model-dependent	SHAP and permutation-importance tables for each champion.
Champion identity and full-data artifact provenance	Selection records, model cards, feature manifests, and training-data fingerprints.
No paper leakage and nested tuning contract	Split-assignment checks and evaluation integrity records.

A.5 Oral-defense reconciliation

The following concise answers reconcile statements that can otherwise appear contradictory.

Why can the model be called strong if grouped performance is moderate? It is strong for interpolation and useful but limited for transfer. The report does not collapse those questions into one score.

Why are Rivera/Jose results above 0.9 while this grouped FSR R^2 is 0.458? Their headline results concern a narrower wire database under random ten-fold cross-validation; this study’s 0.458 holds out entire papers. The protocols, source composition, features, units, and algorithms are not equivalent [9].

Does paper-grouped validation prove performance on a future spacecraft? No. It is stronger internal evidence against paper leakage, but future conditions may differ by laboratory, material, facility, gravity regime, and mission environment. Prospective external validation remains necessary.

Why use XGBoost if random forest wins FSR? XGBoost is a strong tabular prior and wins ignition, consistent with the literature [11]; model selection is empirical and task-specific. Random forest offered the preferred grouped FSR result in the tested portfolio.

Does SHAP identify combustion mechanisms? No. SHAP explains the fitted prediction relative to a background distribution. Correlation, derived features, missingness, and source confounding preclude causal interpretation.

Did bootstrap augmentation create new experiments? No. It resampled observed training rows, inherited their paper IDs, and changed weighting. It added no new physical support.

Can a low predicted risk certify a material? No. The artifacts support screening and prioritisation only. Qualification, certification, and expert hazard assessment remain authoritative.

Why not deploy a PINN immediately? The database mainly contains heterogeneous tabular endpoints, not common-domain fields with agreed governing equations and boundary conditions. Near-term physically constrained or hybrid models are more defensible [8].

A.6 Model portfolios

The executable YAML defines the configured portfolio: 34 FSR candidates (10 XGBoost, seven random forest, six KNN, five MLP, and six SVR) and 26 ignition candidates. The persisted evidence is narrower: the merged FSR manifest contains 25 candidates, each with real-data (**rd**) and bootstrap-augmented (**rd_bt**) conditions, while the formal ignition selection record omits MLP despite available shard metrics. The selected champions

Table A.4: Configured model portfolios.

Task	Family	Candidate hypotheses
FSR	XGBoost (10)	Baseline; square-root/inverse paper weights; absolute-error and pseudo-Huber objectives; signed-log($1 + y$) target; paper bagging.
FSR	Random forest (7)	Baseline and paper weights; absolute-error criteria; signed-log target variants.
FSR	KNN (6)	Uniform or distance neighbours, each with no, square-root, or inverse paper weighting.
FSR	MLP (5)	Baseline; paper weighting; signed-log target variants.
FSR	SVR (6)	Linear or RBF kernels, each with no, square-root, or inverse paper weighting.
Ignition	XGBoost (8)	Baseline; paper-only and class-only weights; combined square-root/inverse weights; focal-loss variants.
Ignition	Decision tree (3)	Baseline and combined class plus square-root/inverse paper weighting.
Ignition	KNN (5)	Baseline; weighted variants; distance-neighbour variants.
Ignition	MLP (5)	Baseline; weighted variants; focal-loss variants.
Ignition	SVM (5)	RBF baseline/weighted variants and linear-kernel weighted variants.

are `rf_baseline [rd]` for both FSR protocols, `xgb_paper_sqrt` for ignition interpolation, and `xgb_baseline` for ignition extrapolation. These names identify frozen configurations within the persisted comparison pools and do not imply universal superiority outside the recorded schema and validation design.

Bibliography

- [1] Tianqi Chen and Carlos Guestrin. XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 785–794, 2016. doi: 10.1145/2939672.2939785.
- [2] Augustin Guibaud, Guillaume Legros, Jean-Louis Consalvi, and Jose Torero. Fire safety in spacecraft: Past incidents and deep space challenges. *Acta Astronautica*, 195:344–354, 2022. doi: 10.1016/j.actaastro.2022.01.021.
- [3] Michael C. Johnston, James S. T’ien, Derek E. Muff, Xiaoyang Zhao, Sandra L. Olson, and Paul V. Ferkul. Self induced buoyant blow off in upward flame spread on thin solid fuels. *Fire Safety Journal*, 71:279–286, 2015. doi: 10.1016/j.firesaf.2014.11.007.
- [4] George Em Karniadakis, Ioannis G. Kevrekidis, Lu Lu, Paris Perdikaris, Sifan Wang, and Liu Yang. Physics-informed machine learning. *Nature Reviews Physics*, 3:422–440, 2021. doi: 10.1038/s42254-021-00314-5.
- [5] Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, pages 2980–2988, 2017.
- [6] NASA. Apollo 1 prime crew portrait. NASA photograph, 1967. Public-domain U.S. Government work.
- [7] NASA. Apollo 204 command module during the post-fire investigation. NASA photograph, 1967. Public-domain U.S. Government work.
- [8] Maziar Raissi, Paris Perdikaris, and George Em Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, 2019. doi: 10.1016/j.jcp.2018.10.045.
- [9] Jose Rivera, Daniel San Martin, Carlos Fernandez-Pello, Michael J. Gollner, Augustin Guibaud, Sandra Olson, and Dennis Stocker. Using data categorization and augmentation strategies to improve machine learning frameworks for flame spread over electrical wires. Manuscript supplied with the project literature corpus; publication year and venue are not stated in the repository copy, n.d.
- [10] Ankit Rohatgi. WebPlotDigitizer. <https://automeris.io/WebPlotDigitizer/>, n.d. Browser-based graph digitization software.
- [11] Ravid Shwartz-Ziv and Amitai Armon. Tabular data: Deep learning is not all you need. *Information Fusion*, 81:84–90, 2022. doi: 10.1016/j.inffus.2021.11.011.